

# OpenMP


# OpenMP: A little history

- Standard introduced in 1996
  - An industry standard after multiple companies and extended their own compilers to virtually the same tasks
  - Most of the predecessors offered more flexibility than the OpenMP standard
- Extension built into most commercial Fortran90, C, and C++ compilers

# OpenMP

- The basics
- Loops
- Parallel environment
- Examples

# OpenMP program diagram



# Hello world

```
program hello_world  
write(0,*) "Hello world"  
end program
```

```
>ifort hello.f90  
>./a.out  
Hello world
```

```
#include<stdio.h>  
int main (int argc, char **argv)  
{  
    fprintf(stderr,"Hello world");  
}
```

```
>icc hello.c  
>./a.out  
Hello world
```

# Hello world

```
program hello_world
!$OMP PARALLEL
write(0,*) "Hello world"
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world
Hello world
Hello world
```

```
#include<stdio.h>
int main (int argc, char **argv)
{
    #pragma omp parallel
    fprintf(stderr,"Hello world");
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world
Hello world
Hello world
```

# Hello world

```
program hello_world
!$OMP PARALLEL
write(0,*) "Hello world"
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world
Hello world
Hello world
```

```
#include<stdio.h>
int main (int argc, char **argv)
{
#pragma omp parallel
    fprintf(stderr,"Hello world");
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world
Hello world
Hello world
```

Begin and end a parallel environment in  
Fortran90

# Hello world

```
program hello_world
!$OMP PARALLEL
write(0,*) "Hello world"
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world
Hello world
Hello world
```

```
#include<stdio.h>
int main (int argc, char **argv)
{
    #pragma omp parallel
    fprintf(stderr,"Hello world");
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world
Hello world
Hello world
```

Delineate a parallel region in C



# Hello world

```
program hello_world
!$OMP PARALLEL
write(0,*) "Hello world"
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world
Hello world
Hello world
```

```
#include<stdio.h>
int main (int argc, char **argv)
{
    #pragma omp parallel
    fprintf(stderr,"Hello world");
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world
Hello world
Hello world
```

Tell how many different copies of the programs to run simultaneously

# Hello world

```
program hello_world
!$OMP PARALLEL
write(0,*) "Hello world"
!$OMP END PARALLEL
end program
```

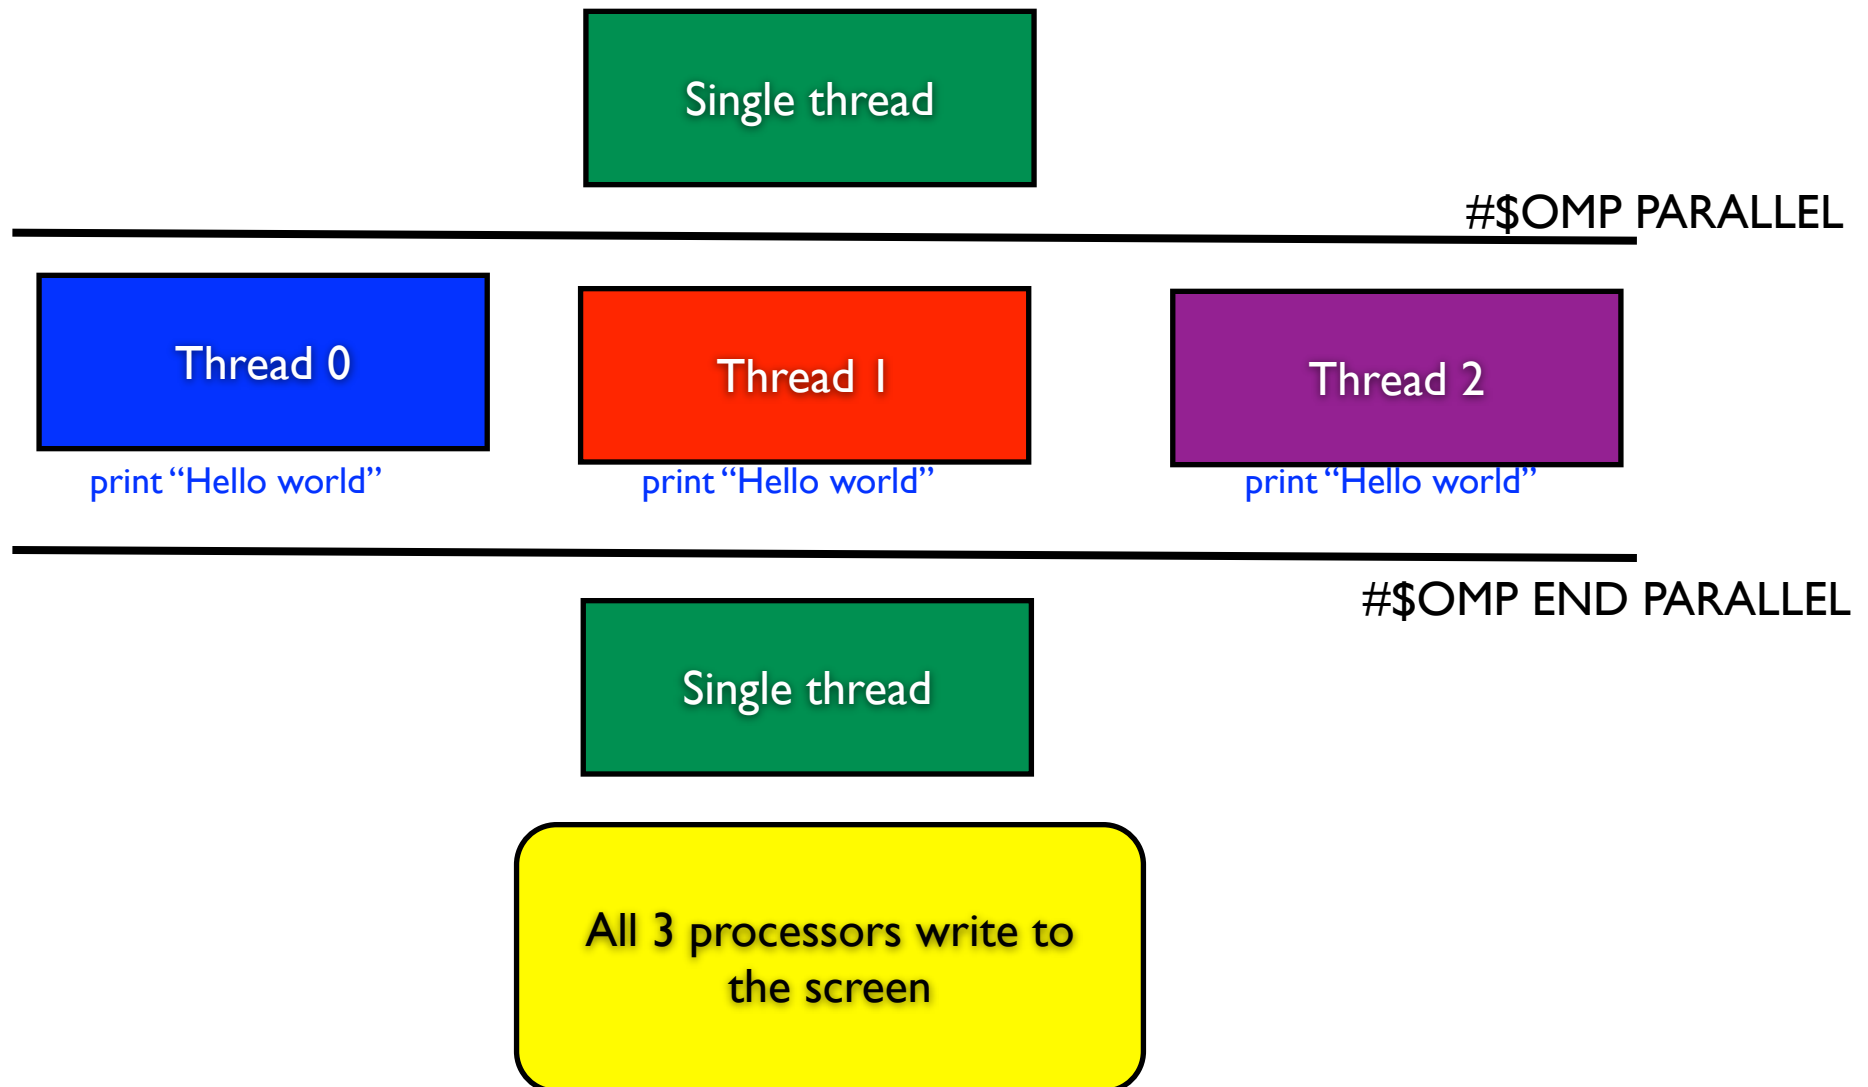
```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world
Hello world
Hello world
```

```
#include<stdio.h>
int main (int argc, char **argv)
{
    #pragma omp parallel
    fprintf(stderr,"Hello world");
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world
Hello world
Hello world
```

Compile with openmp extensions

# What is going on



# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
!$OMP PARALLEL
write(0,*) "Hello world",omp_get_thread_num(),"of",&
  omp_get_num_threads()
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world  1 of  3
Hello world  2 of  3
Hello world  0 of  3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
  #pragma omp parallel
  fprintf(stderr,"Hello world %d of %d\n",
    omp_get_thread_num(), omp_get_num_threads());
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
!$OMP PARALLEL
write(0,*) "Hello world",omp_get_thread_num(),"of",&
  omp_get_num_threads()
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world  1 of  3
Hello world  2 of  3
Hello world  0 of  3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
  #pragama omp parallel
  fprintf(stderr,"Hello world %d of %d\n",
    omp_get_thread_num(), omp_get_num_threads());
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

Note external functions in Fortran90

# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
!$OMP PARALLEL
write(0,*) "Hello world",omp_get_thread_num(),"of",&
  omp_get_num_threads()
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world  1 of  3
Hello world  2 of  3
Hello world  0 of  3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
  #pragma omp parallel
  fprintf(stderr,"Hello world %d of %d\n",
    omp_get_thread_num(), omp_get_num_threads());
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

Function prototypes definitions

# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
!$OMP PARALLEL
write(0,*) "Hello world",omp_get_thread_num(),"of",&
  omp_get_num_threads()
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world 1 of 3
Hello world 2 of 3
Hello world 0 of 3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
  #pragma omp parallel
  fprintf(stderr,"Hello world %d of %d\n",
    omp_get_thread_num(), omp_get_num_threads());
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

Return what thread the I am.

# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
!$OMP PARALLEL
write(0,*) "Hello world",omp_get_thread_num(),"of",&
  omp_get_num_threads()
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world 1 of 3
Hello world 2 of 3
Hello world 0 of 3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
  #pragma omp parallel
  fprintf(stderr,"Hello world %d of %d\n",
    omp_get_thread_num(), omp_get_num_threads());
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

The number of threads in the parallel environment.



# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
!$OMP PARALLEL
write(0,*) "Hello world",omp_get_thread_num(),"of",&
  omp_get_num_threads()
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world  1 of  3
Hello world  2 of  3
Hello world  0 of  3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
  #pragama omp parallel
  fprintf(stderr,"Hello world %d of %d\n",
    omp_get_thread_num(), omp_get_num_threads());
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

First thread is #0

# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
!$OMP PARALLEL
write(0,*) "Hello world",omp_get_thread_num(),"of",&
  omp_get_num_threads()
!$OMP END PARALLEL
end program
```

```
>setenv OMP_NUM_THREADS 3
>ifort -openmp hello.f90
>./a.out
Hello world 1 of 3
Hello world 2 of 3
Hello world 0 of 3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
  #pragma omp parallel
  fprintf(stderr,"Hello world %d of %d\n",
    omp_get_thread_num(), omp_get_num_threads());
}
```

```
>setenv OMP_NUM_THREADS 3
>icc -openmp hello.c
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

Print statements not necessarily in order  
or consistent.

# Hello world

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
    int impi,nmpi;
    #pragama omp parallel
    {
        ith=omp_get_thread_num();
        nth=omp_get_num_threads();
    }
    fprintf(stderr,"Hello world %d of %d\n",
        ith,nth));
}
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
    int impi,nmpi;
    #pragama omp parallel
    {
        ith=omp_get_thread_num();
        nth=omp_get_num_threads();
        fprintf(stderr,"Hello world %d of %d\n",
            ith,nth));
    }
}
```

# Hello world

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
    int impi,nmpi;
    #pragama omp parallel
    {
        ith=omp_get_thread_num();
        nth=omp_get_num_threads();
    }
    fprintf(stderr,"Hello world %d of %d\n",
        ith,nth));
}
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
    int impi,nmpi;
    #pragama omp parallel
    {
        ith=omp_get_thread_num();
        nth=omp_get_num_threads();
        fprintf(stderr,"Hello world %d of %d\n",
            ith,nth));
    }
}
```

# Hello world

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
    int ith,nth;
    #pragma omp parallel
    {
        ith=omp_get_thread_num();
        nth=omp_get_num_threads();
    }
    fprintf(stderr,"Hello world %d of %d\n",
        ith,nth));
}
```

```
>./a.out
Hello world 1 of 3
>./a.out
Hello world 2 of 3
>./a.out
Hello world 1 of 3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
    int ith,nth;
    #pragma omp parallel
    {
        ith=omp_get_thread_num();
        nth=omp_get_num_threads();
    }
    fprintf(stderr,"Hello world %d of %d\n",
        ith,nth));
}
```

Result not necessarily consistent

# Hello world

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
    int impi,nmpi;
    #pragma omp parallel
    {
        ith=omp_get_thread_num();
        nth=omp_get_num_threads();
    }
    fprintf(stderr,"Hello world %d of %d\n",
        ith,nth));
}
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
    int impi,nmpi;
    #pragma omp parallel
    {
        ith=omp_get_thread_num();
        nth=omp_get_num_threads();
        fprintf(stderr,"Hello world %d of %d\n",
            ith,nth));
    }
}
```

```
>./a.out
Hello world 1 of 3
Hello world 2 of 3
Hello world 2 of 3
```

# What is going on

Single thread

#\$OMP PARALLEL

ith=0

print ith (0)

ith=1

print ith (2)

ith=2

print ith(2)

Memory location can be changed  
between setting ith and printing ith

# Public vs private

Within a parallel region every variable is assigned either accessible variable is assigned either a public or private status.



# Public variables

- All threads access the same memory location (transferred to the heap)
- Thread timing is unpredictable so changes to the variable are unpredictable and inconsistent
- With some caveats that we will discuss later
- All things allocated on the stack must be public

# Heap

- func()
  - $a=b+c$

# Private variables

- Each thread has its own copy
- **Private variables exist on the stack**
  - stack requirements increase by nthreads
  - by default they disappear when exiting a parallel region

# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
integer :: ith,nth
!$OMP PARALLEL default(shared) private(ith,nth)
ith=omp_get_thread_num()
nth=omp_get_num_threads()
write(0,*) "Hello world",ith,nth
!$OMP END PARALLEL
end program
```

```
>./a.out
Hello world  1 of  3
Hello world  2 of  3
Hello world  0 of  3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
int ith, nth;
#pragma omp parallel default(shared) private(ith,nth)
{
ith=omp_get_thread_num()
nth=omp_get_num_threads()
fprintf(stderr,"Hello world %d of %d\n",
omp_get_thread_num(), omp_get_num_threads());
}
```

```
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
integer :: ith,nth
!$OMP PARALLEL default(shared) private(ith,nth)
ith=omp_get_thread_num()
nth=omp_get_num_threads()
write(0,*) "Hello world",ith,nth
!$OMP END PARALLEL
end program
```

```
>./a.out
Hello world  1 of  3
Hello world  2 of  3
Hello world  0 of  3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
int ith, nth;
#pragma omp parallel default(shared) private(ith,nth)
    ith=omp_get_thread_num()
    nth=omp_get_num_threads()
    fprintf(stderr,"Hello world %d of %d\n",
        omp_get_thread_num(), omp_get_num_threads());
}
```

```
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

By default all variables are shared

# Hello world

```
program hello_world
integer, external :: omp_get_thread_num
integer, external :: omp_get_num_threads
integer :: ith,nth
!$OMP PARALLEL default(shared) private(ith,nth)
ith=omp_get_thread_num()
nth=omp_get_num_threads()
write(0,*) "Hello world",ith,nth
!$OMP END PARALLEL
end program
```

```
>./a.out
Hello world  1 of  3
Hello world  2 of  3
Hello world  0 of  3
```

```
#include<stdio.h>
#include<omp.h>
int main (int argc, char **argv)
{
int ith, nth;
#pragma omp parallel default(shared) private(ith,nth)
    ith=omp_get_thread_num()
    nth=omp_get_num_threads()
    fprintf(stderr,"Hello world %d of %d\n",
        omp_get_thread_num(), omp_get_num_threads());
}
```

```
>./a.out
Hello world 0 of 3
Hello world 2 of 3
Hello world 1 of 3
```

The list of variables that are private

# OMP do/for construct

- A loop is executed by multiple threads
- Data parallel approach

# Vector multiplication

```
subroutine vec_mult(out,in1,in2)
  real :: out(:),in1(:),in2(:)
  integer :: i
!$OMP PARALLEL default(shared) private(i)
!$OMP DO
do i=1,size(out)
  out(i)=in1(i)*in2(i)
end do
!$OMP END DO
!$OMP END PARALLEL
end subroutine
```

```
void vec_mult(int n1,float *out,float *in1,float *in2)
{
  int i;
  #pragma omp parallel default(shared) private(i)
  #pragam omp for
  for(i=0; i < n; i++) out[i]=in1[i]*in2[i];
}
```



# Vector multiplication

```
subroutine vec_mult(out,in1,in2)
  real :: out(:),in1(:),in2(:)
  integer :: i
  !$OMP PARALLEL default(shared) private(i)
  !$OMP DO
  do i=1,size(out)
    out(i)=in1(i)*in2(i)
  end do
  !$OMP END DO
  !$OMP END PARALLEL
end subroutine
```

```
void vec_mult(int n1,float *out,float *in1,float *in2)
{
  int i;
  #pragma omp parallel default(shared) private(i)
  #pragam omp for
  for(i=0; i < n; i++) out[i]=in1[i]*in2[i];
}
```

Parallelize loop

# Vector multiplication

```
subroutine vec_mult(out,in1,in2)
  real :: out(:),in1(:),in2(:)
  integer :: i
  !$OMP PARALLEL DO default(shared) private(i)
  do i=1,size(out)
    out(i)=in1(i)*in2(i)
  end do
  !$OMP END PARALLEL DO
end subroutine
```

```
void vec_mult(int n1,float *out,float *in1,float *in2)
{
  int i;
  #pragma omp parallel for default(shared) private(i)
  for(i=0; i < n; i++) out[i]=in1[i]*in2[i];
}
```

Loop and parallel compiler directives can be combined.

# Dot product (version 1)

```
double precision function dot_product(in1,in2)
  real :: in1(:),in2(:)
  integer :: i
  dot_product=0
!$OMP PARALLEL DO default(shared) private(i)
do i=1,size(in1)
  dot_product=dot_product+in1(i)*in2(i)
end do
!$OMP END PARALLEL DO
end function
```

```
double dot_product(int n, float *in1,*in2)
{
  int i;
  double out=0.;
  #pragma omp parallel for default(shared) private(i)
  for(i=0; i < n; i++)out+=in1[i]*in2[i];
  return(out);
}
```

# Dot product (version 1)

```
double precision function dot_product(in1,in2)
  real :: in1(:),in2(:)
  integer :: i
  dot_product=0
!$OMP PARALLEL DO default(shared) private(i)
do i=1,size(in1)
  dot_product=dot_product+in1(i)*in2(i)
end do
!$OMP END PARALLEL DO
end function
```

```
double dot_product(int n, float *in1,*in2)
{
  int i;
  double out=0.;
  #pragma omp parallel for default(shared) private(i)
  for(i=0; i < n; i++)out+=in1[i]*in2[i];
  return(out);
}
```

Race condition (different threads may not send to heap updated out quickly enough)

# Dot product (version 1)

```
double precision function dot_product(in1,in2)
  real :: in1(:),in2(:)
  integer :: i
  dot_product=0
  !$OMP PARALLEL DO default(shared) private(i) &
do i=1,size(in1)
  !$OMP ATOMIC
  dot_product=dot_product+in1(i)*in2(i)
end do
!$OMP END PARALLEL DO
end function
```

```
double dot_product(int n, float *in1,*in2)
{
  int i;
  double out=0.;
  #pragma omp parallel for default(shared) private(i)
  {
    for(i=0; i < n; i++){
      #pragma omp atomic
      out+=in1[i]*in2[i];
    }
  }
  return(out);
}
```

Guarantees that statement will be completed and memory location updated before another processors accesses it.  
Limited number of constructs (+, -, etc.)


# Dot product (version 1)

```
double precision function dot_product(in1,in2)
  real :: in1(:),in2(:)
  integer :: i
  dot_product=0
  !$OMP PARALLEL DO default(shared) private(i)
  do i=1,size(in1)
    !$OMP CRITICAL
    dot_product=dot_product+in1(i)*in2(i)
  !$OMP END CRITICAL
  end do
  !$OMP END CRITICAL
!$OMP END PARALLEL DO
end function
```


```
double dot_product(int n, float *in1,*in2)
{
  int i;
  double out=0.;
  #pragma omp parallel for default(shared) private(i)
  {
    for(i=0; i < n; i++){
      #pragma omp critical
      out+=in1[i]*in2[i];
    }
  }
  return(out);
}
```

Guarantees that only one thread at a time will enter the code section.

# Theoretical speedup




# Compute time

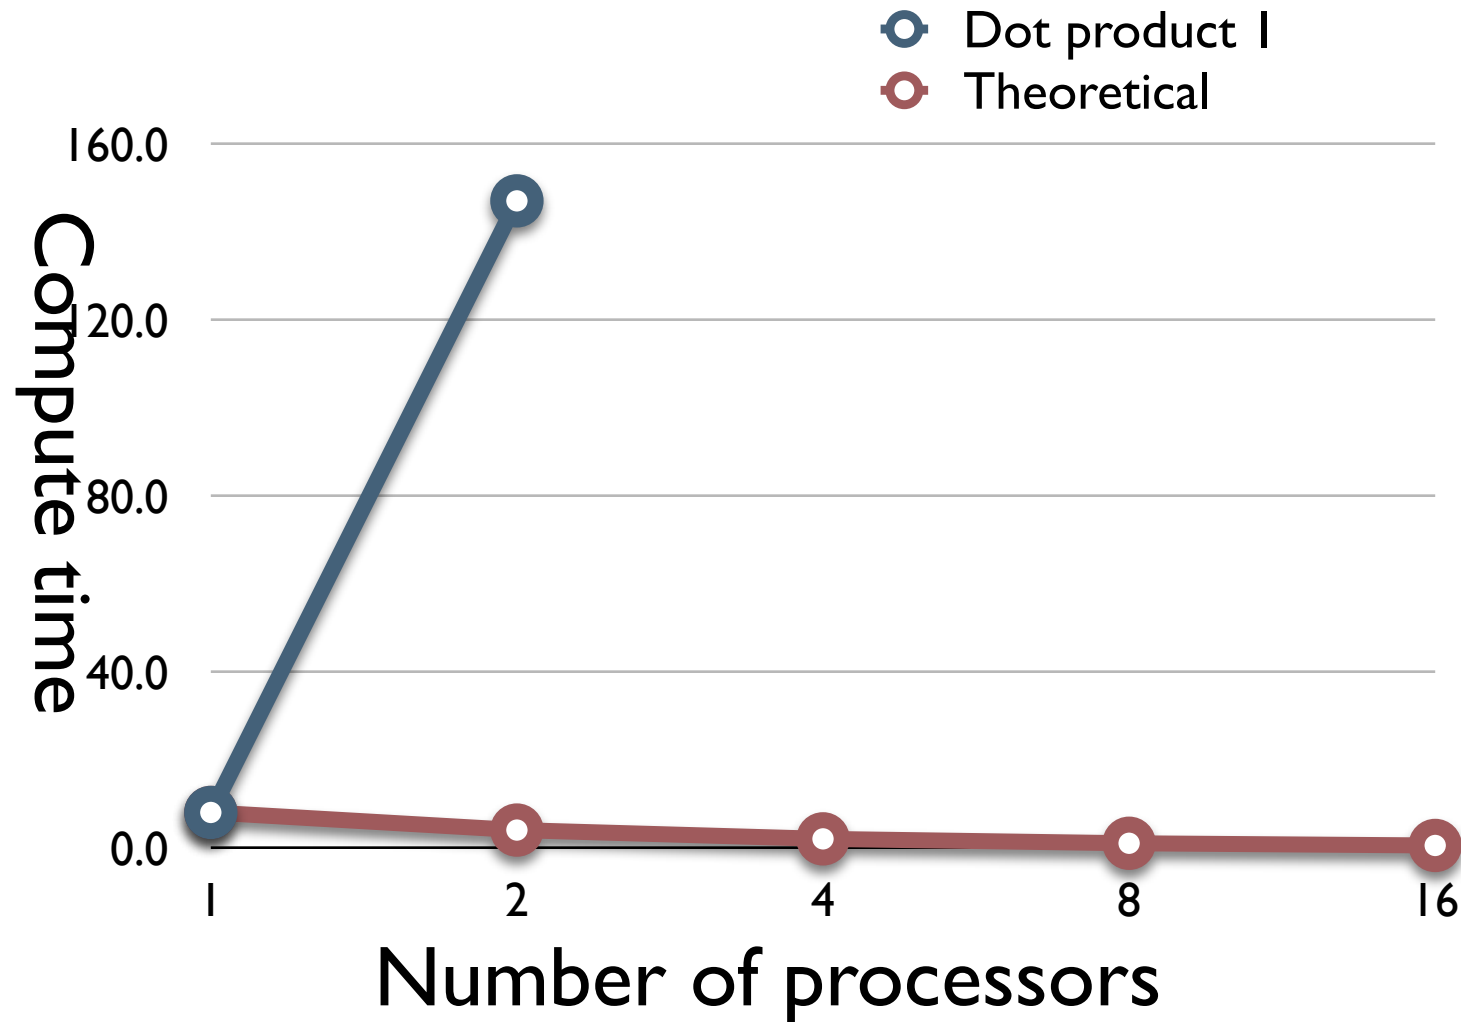





# Dot product(I)



# Dot product(I)



# Dot product(I)



# What is going on

```
out+=in l [0]*in2[0]
```

# First thread creates block

```
out+=in1[0]*in2[0]
```

Block

# Other thread hits block



# Other thread check

status

out+=in1[0]\*in2[0]

out+=in1[1]\*in2[1]

out+=in1[2]\*in2[2]


Block

# Other threads sleep






# Next thread is unblocked

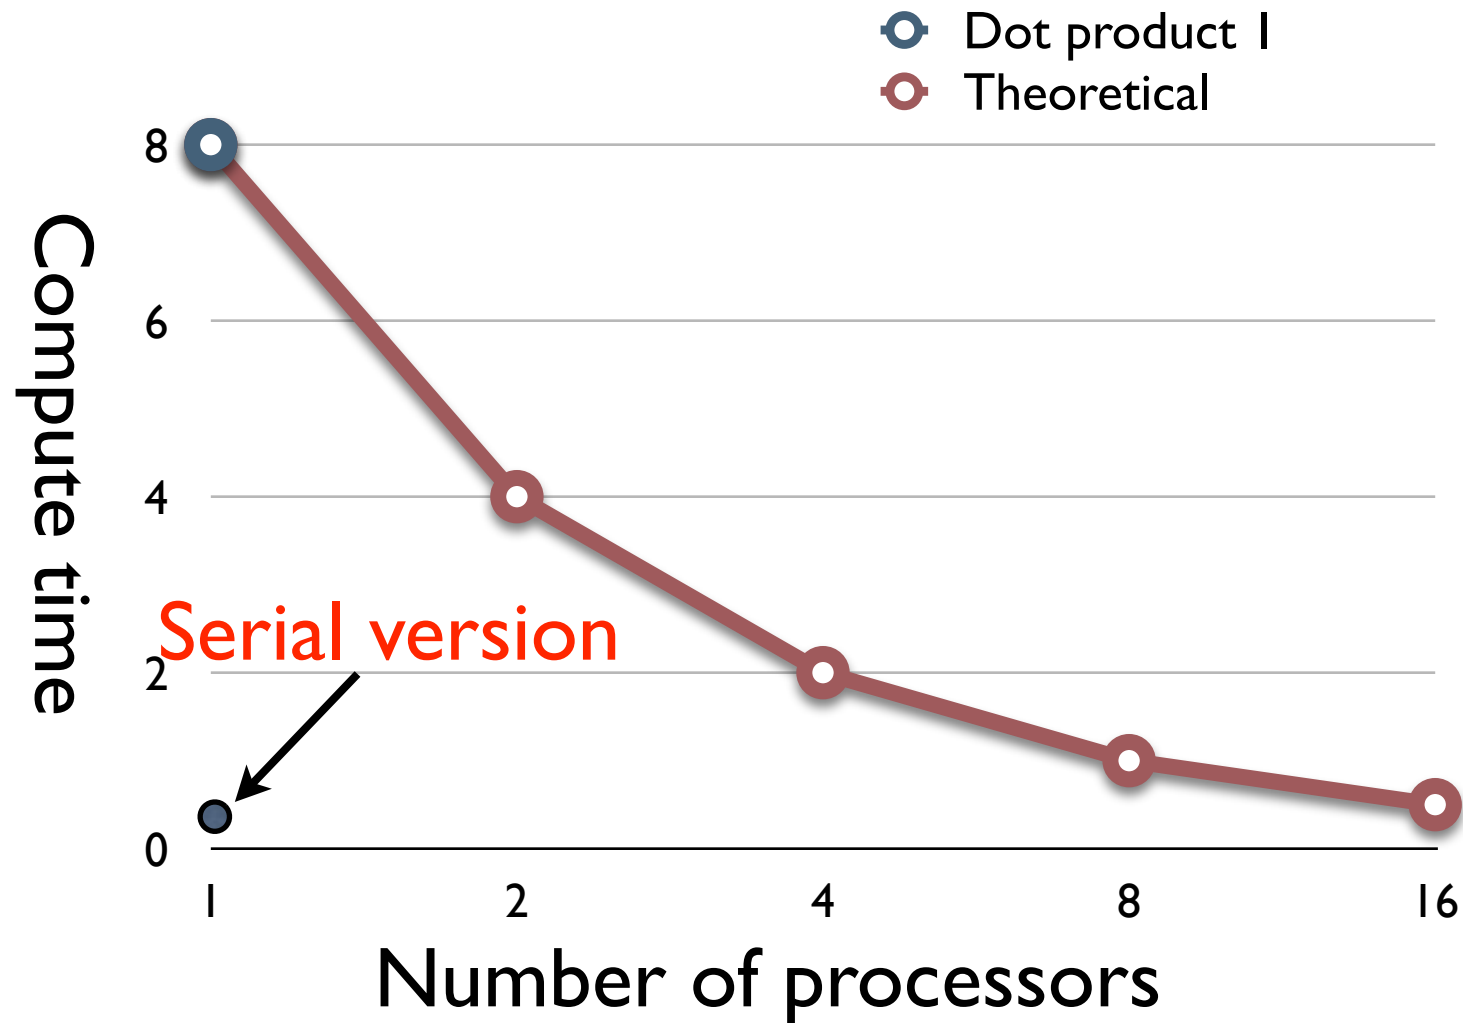


# Next thread hits block




Lock checking/sleep combo is a non-trivial cost

# Dot product(I)



# Next thread hits block



The cost of setting up the locks 20x the cost of doing the addition

# Dot product (version 2)

```
double precision function dot_product(in1,in2)
  real :: in1 (:),in2(:)
  double precision :: dots(100)
  integer :: i,ith
  integer,external :: omp_get_thread_num
  dots=0;dot_product=0
!$OMP PARALLEL default(shared) private(ith,i)
ith=omp_get_thread_num()
!$OMP DO
do i=1,size(in1)
  dots(ith+1)=dots(ith+1)+in1(i)*in2(i)
end do
!$OMP END DO
!$OMP END PARALLEL
dot_product=sum(dots)
end function
```

```
double dot_product(int n, float *in1, *in2)
{
  int i,ith;
  double dots[100], out=0.;
  for(i=0; i < 100; i++) dots[i]=0.;
  #pragma omp parallel default(shared) private(ith,i)
  ith=omp_get_thread_num();
  #pragma omp for
  {
    for(i=0; i < n; i++){
      #pragma omp critical
      dots[ith]+=in1[i]*in2[i];
    }
  }
  for(i=0; i < 100; i++) out+=dots[i];
  return(out);
}
```

# Dot product (version 2)

```
double precision function dot_product(in1,in2)
  real :: in1 (:),in2(:)
  double precision :: dots(100)
  integer :: i,ith
  integer,external :: omp_get_thread_num
  dots=0;dot_product=0
  !$OMP PARALLEL default(shared) private(i,ith)
  ith=omp_get_thread_num()
  !$OMP DO
  do i=1,size(in1)
    dots(ith+1)=dots(ith+1)+in1(i)*in2(i)
  end do
  !$OMP END DO
  !$OMP END PARALLEL
  dot_product=sum(dots)
end function
```

```
double dot_product(int n, float *in1,*in2)
{
  int i,ith;
  double dots[100], out=0.;
  for(i=0; i < 100; i++) dots[i]=0.;
  #pragma omp parallel default(shared) private(ith,i)
  ith=omp_get_thread_num();
  #pragma omp for
  {
    for(i=0; i < n; i++){
      dots[ith]+=in1[i]*in2[i];
    }
  }
  for(i=0; i < 100; i++) out+=dots[i];
  return(out);
}
```

Grab thread number

# Dot product (version 2)

```
double precision function dot_product(in1,in2)
  real :: in1(:),in2(:)
  double precision :: dots(100)
  integer :: i,ith
  integer,external :: omp_get_thread_num
  dots=0;dot_product=0
  !$OMP PARALLEL default(shared) private(i,ith)
  ith=omp_get_thread_num()
  !$OMP DO
  do i=1,size(in1)
    dots(ith+1)=dots(ith+1)+in1(i)*in2(i)
  end do
  !$OMP END DO
  !$OMP END PARALLEL
  dot_product=sum(dots)
end function
```

```
double dot_product(int n, float *in1,*in2)
{
  int i,ith;
  double dots[100], out=0.;
  for(i=0; i < 100; i++) dots[i]=0.;
  #pragma omp parallel default(shared) private(ith,i)
  ith=omp_get_thread_num();
  #pragma omp for
  {
    for(i=0; i < n; i++){
      dots[ith]+=in1[i]*in2[i];
    }
  }
  for(i=0; i < 100; i++) out+=dots[i];
  return(out);
}
```

Local dot product sum for each thread

# Dot product (version 2)

```
double precision function dot_product(in1,in2)
  real :: in1(:),in2(:)
  double precision :: dots(100)
  integer :: i,ith
  integer,external :: omp_get_thread_num
  dots=0;dot_product=0
!$OMP PARALLEL default(shared) private(ith,i)
  ith=omp_get_thread_num()
!$OMP DO
  do i=1,size(in1)
    dots(ith+1)=dots(ith+1)+in1(i)*in2(i)
  end do
!$OMP END DO
!$OMP END PARALLEL
dot_product=sum(dots)
end function
```

```
double dot_product(int n, float *in1,*in2)
{
  int i,ith;
  double dots[100], out=0.;
  for(i=0; i < 100; i++) dots[i]=0.;
  #pragma omp parallel default(shared) private(ith,i)
  ith=omp_get_thread_num();
  #pragma omp for
  {
    for(i=0; i < n; i++){
      dots[ith]+=in1[i]*in2[i];
    }
  }
  for(i=0; i < 100; i++) out+=dots[i];
  return(out);
}
```

Combine local dot product results.



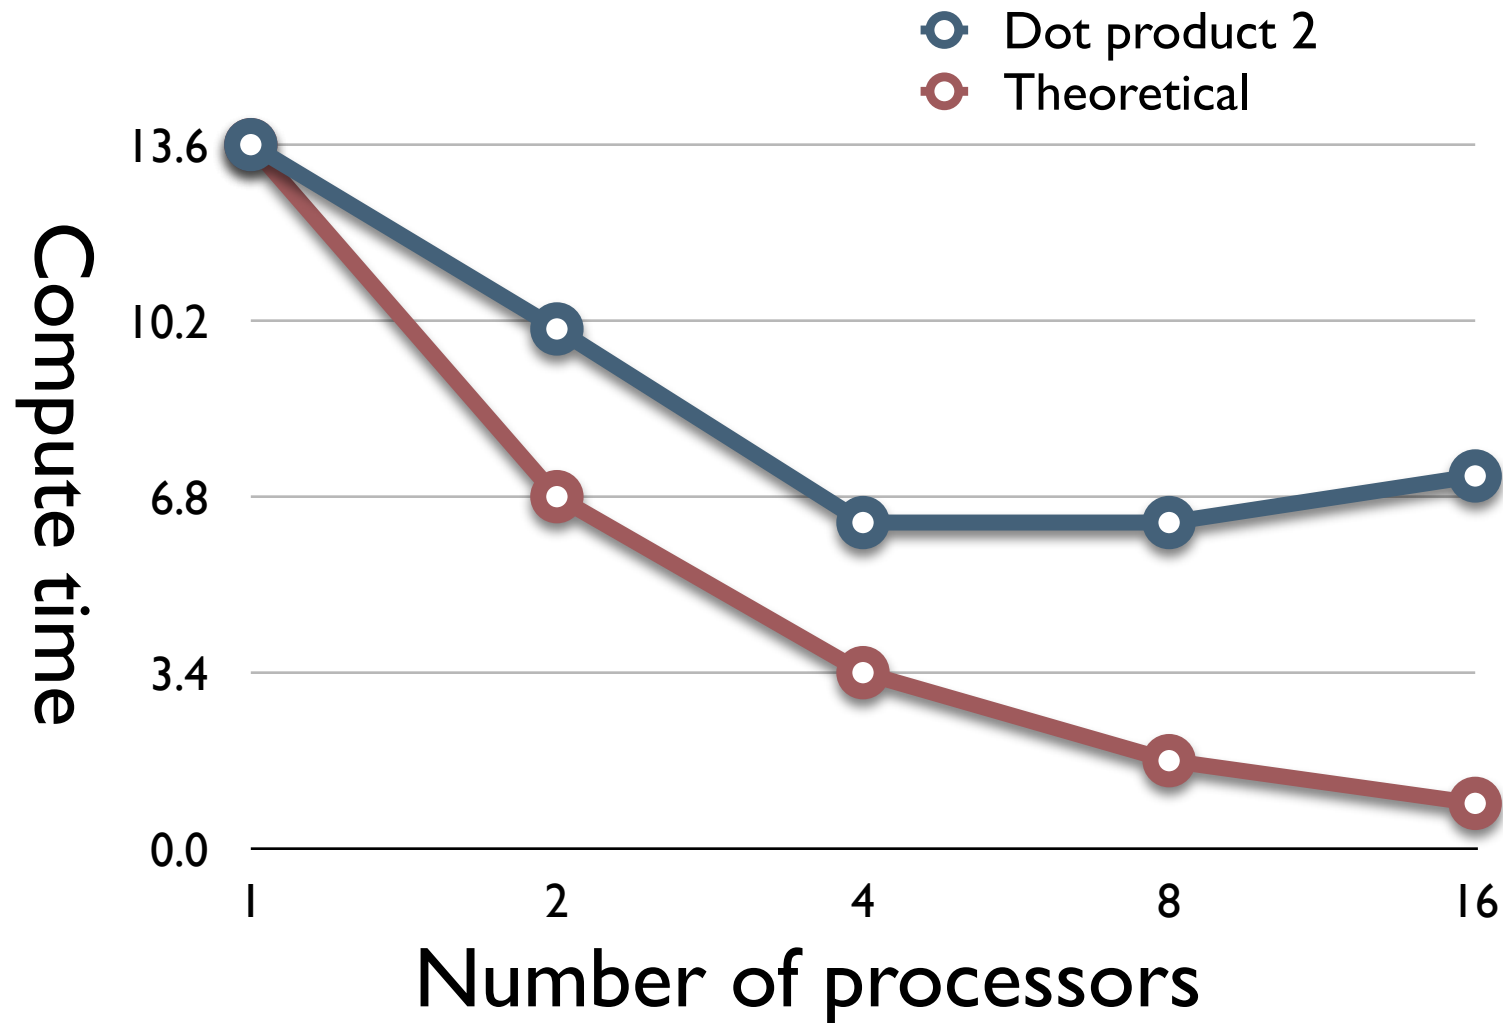
# Dot product (version 2)

```
double precision function dot_product(in1,in2)
  real :: in1 (:),in2(:)
  double precision :: dots(100)
  integer :: i,ith
  integer,external :: omp_get_thread_num
  dots=0;dot_product=0
!$OMP PARALLEL default(shared) private(ith,i)
ith=omp_get_thread_num()
!$OMP DO
do i=1,size(in1)
  dots(ith+1)=dots(ith+1)+in1(i)*in2(i)
end do
!$OMP END DO
!$OMP END PARALLEL
dot_product=sum(dots)
end function
```

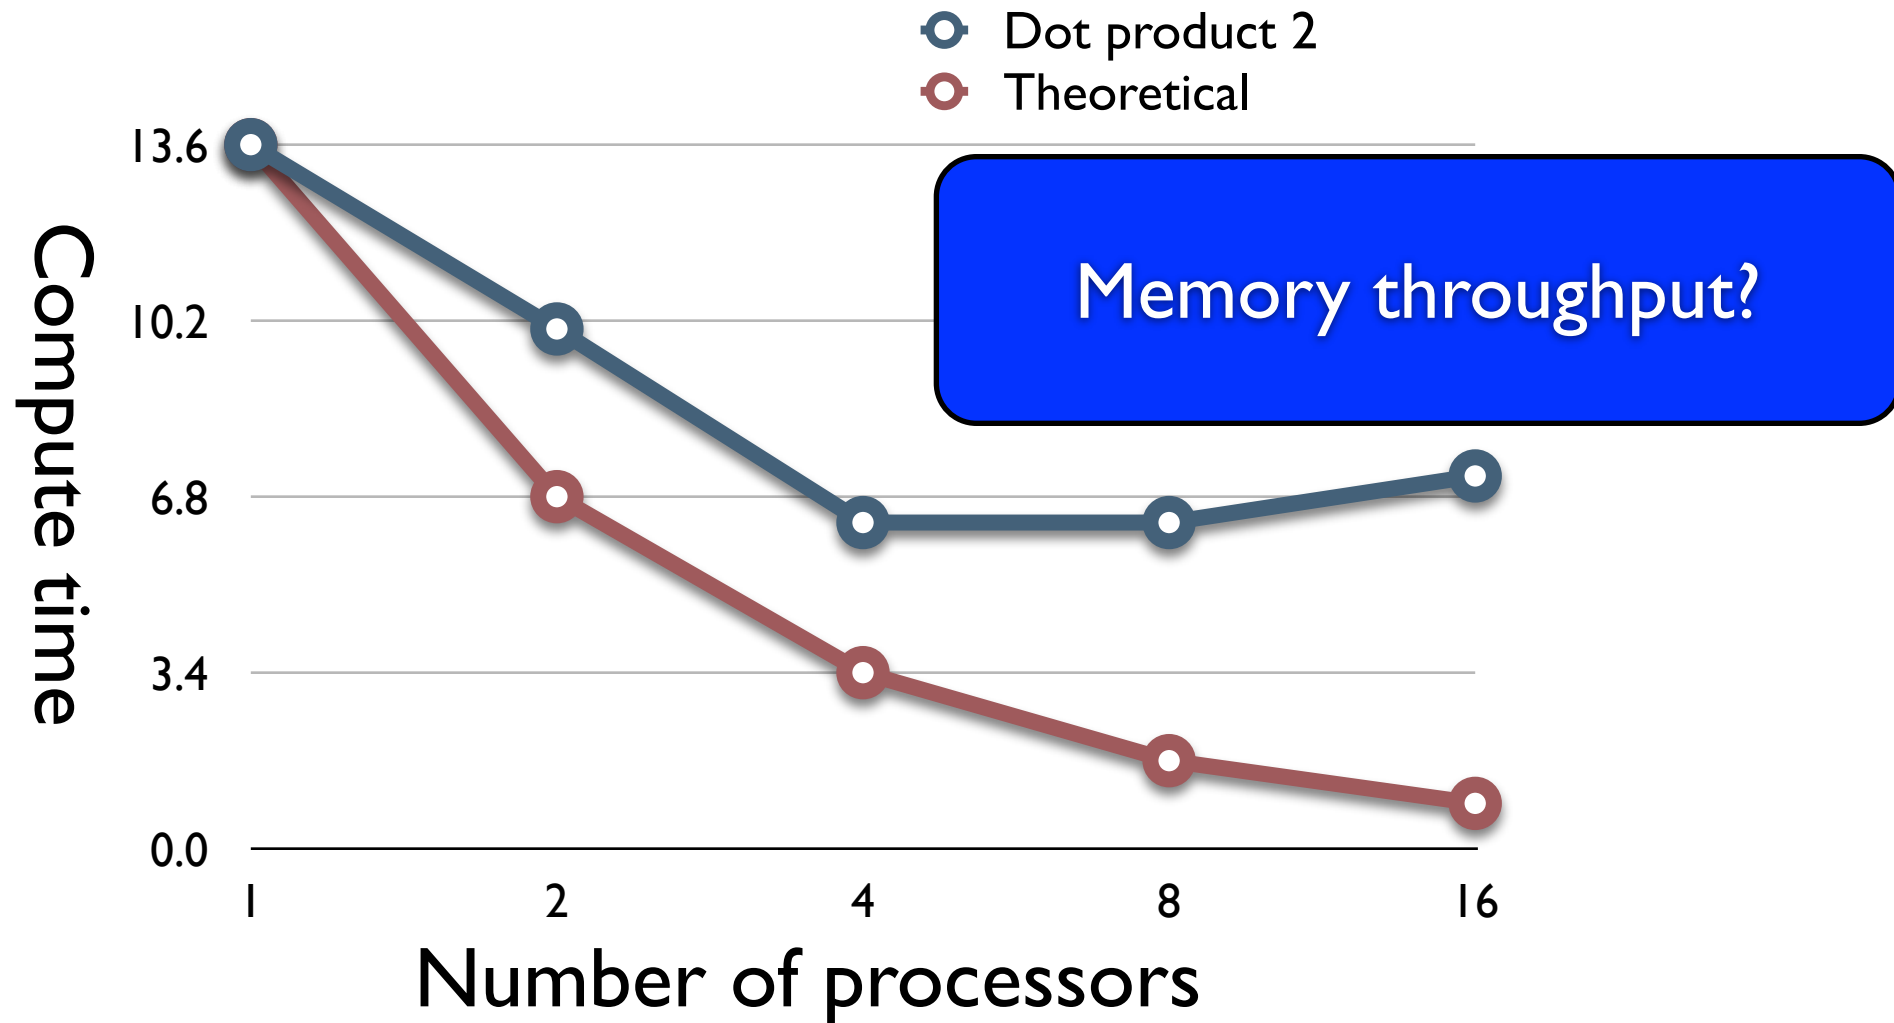
```
double dot_product(int n, float *in1,*in2)
{
  int i,ith;
  double dots[100], out=0.;
  for(i=0; i < 100; i++) dots[i]=0.;
  #pragma omp parallel default(shared) private(ith,i)
  ith=omp_get_thread_num();
  #pragma omp for
  {
    for(i=0; i < n; i++){
      dots[ith]+=in1[i]*in2[i];
    }
  }
  for(i=0; i < 100; i++) out+=dots[i];
  return(out);
}
```

Remove all blocking from the loop.

# Dot product(2)



# Dot product(2)




# Vector distance

```
double precision function vec_dist(in1,in2)
  real :: in1 (:),in2(:)
  double precision :: dots(100)
  integer :: i,ith
  integer,external :: omp_get_thread_num
  dots=0;dot_product=0
!$OMP PARALLEL default(shared) private(ith,i)
ith=omp_get_thread_num()
!$OMP DO
do i=1,size(in1)
  dots(ith+1)=dots(ith+1)+sqrt(in1(i)**2+in2(i)**2)
end do
!$OMP END DO
!$OMP END PARALLEL
vec_dist=sum(dots)
end function
```

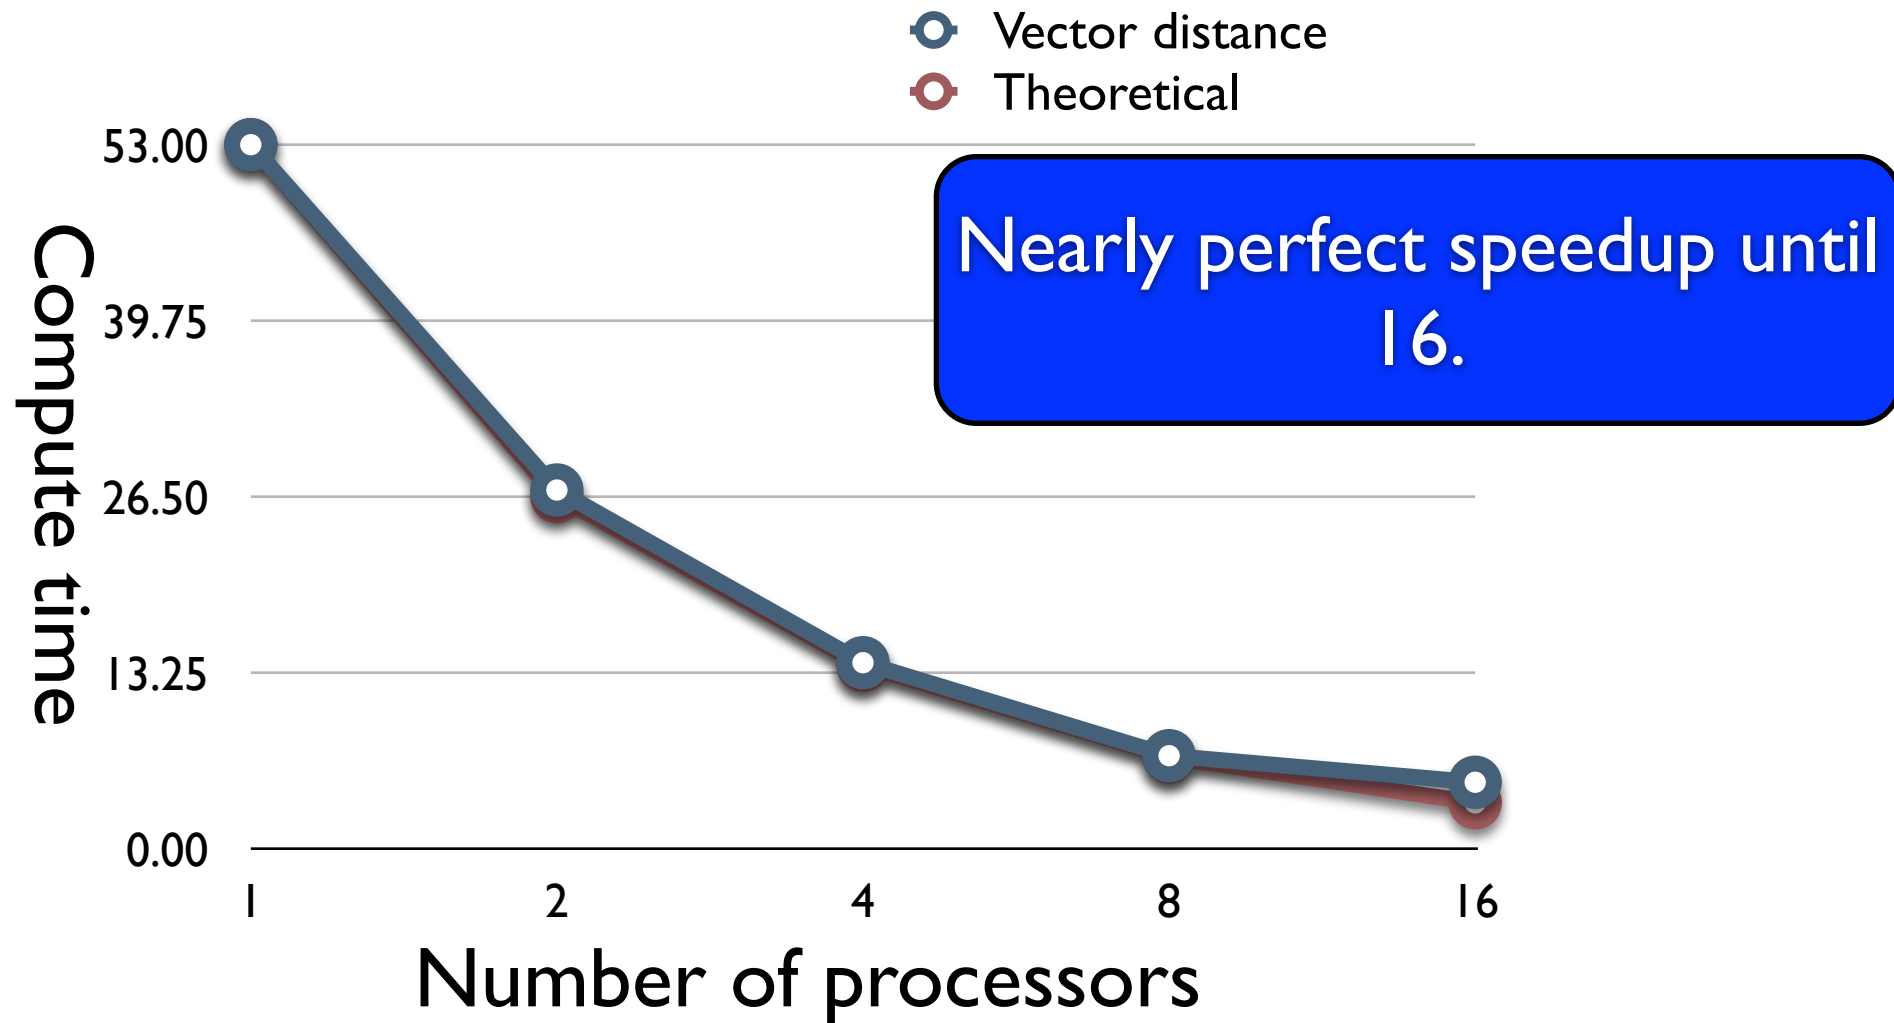
```
double vec_dist(int n, float *in1, *in2)
{
  int i,ith;
  double dots[100], out=0.;
  for(i=0; i < 100; i++) dots[i]=0.;
  #pragma omp parallel default(shared) private(ith,i)
  ith=omp_get_thread_num();
  #pragma omp for
  {
    for(i=0; i < n; i++){
      dots[ith]+=sqrt(in1[i]*in1[i]+in2[i]*in2[i]);
    }
  }
  for(i=0; i < 100; i++) out+=dots[i];
  return(out);
}
```

More operations per datapoint

# Vector distance



# Vector distance



# OpenMP scheduling

- Static - Broken into equal blocks (low overhead)
- Dynamic - Once a thread finishes it is assigned next available loop iteration
- Other scheduling (affinity, guided, runtime)

# Scheduling: Static (2 threads)



**Thread 1**

**Thread 2**



# Scheduling: Dynamic(2 threads)



**Thread 1**

**Thread 2**

# Scheduling: Dynamic(2 threads)



**Thread 1**

**Thread 2**

# Scheduling: Dynamic(2 threads)



**Thread 1**

**Thread 2**

# Scheduling: Dynamic(2 threads)



**Thread 1**

**Thread 2**

# Scheduling: Dynamic(2 threads)



**Thread 1**

**Thread 2**


# Scheduling: Dynamic (2 threads)



**Thread 1**

**Thread 2**


# Scheduling: Dynamic specifying a block size (2 threads)



**Thread 1**

**Thread 2**

# Scheduling: Dynamic specifying a block size (2 threads)



**Thread 1**

**Thread 2**



# Scheduling: Dynamic specifying a block size (2 threads)



**Thread 1**

**Thread 2**

# Scheduling: Dynamic specifying a block size (2 threads)



**Thread 1**

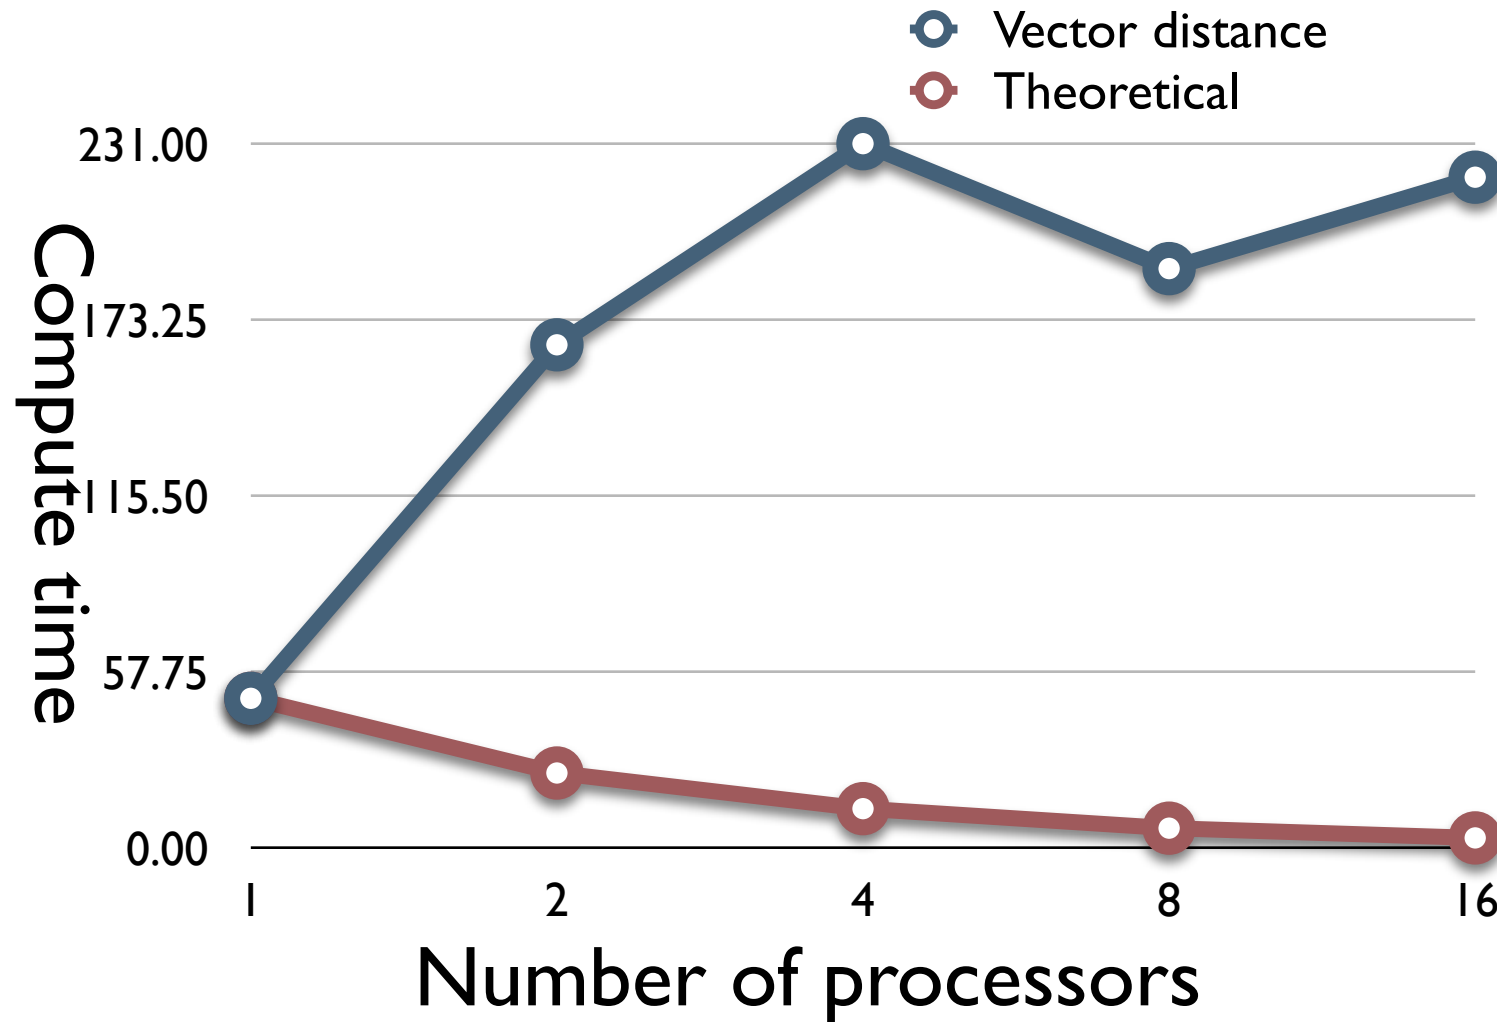
**Thread 2**

# Vector distance

```
double precision function vec_dist(in1,in2)
  real :: in1 (:),in2(:)
  double precision :: dots(100)
  integer :: i,ith
  integer,external :: omp_get_thread_num
  dots=0;dot_product=0
!$OMP PARALLEL default(shared) private(i,ith)
ith=omp_get_thread_num()
!$OMP DO schedule(dynamic,5)
do i=1,size(in1)
  dots(ith+1)=dots(ith+1)+sqrt(in1(i)**2+in2(i)**2)
end do
!$OMP END DO
!$OMP END PARALLEL
vec_dist=sum(dots)
end function
```

```
double vec_dist(int n, float *in1, *in2)
{
  int i,ith;
  double dots[100], out=0.;
  for(i=0; i < 100; i++) dots[i]=0.;
  #pragma omp parallel default(shared) private(ith,i)
  ith=omp_get_thread_num();
  #pragma omp for schedule(dynamic,5)
  {
    for(i=0; i < n; i++){
      dots[ith]+=sqrt(in1[i]*in1[i]+in2[i]*in2[i]);
    }
  }
  for(i=0; i < 100; i++) out+=dots[i];
  return(out);
}
```

# Vector distance: dynamic



# Dot product (version 3)

```
double precision function dot_product(in1,in2)
  real :: in1 (:),in2(:)
  integer :: i,ith
  integer,external :: omp_get_thread_num
  dots=0;dot_product=0
!$OMP PARALLEL default(shared) private(i,ith)
  ith=omp_get_thread_num()
!$OMP REDUCTION(+:dot_product)
!$OMP DO
  do i=1,size(in1)
    dot_product=dots(ith+1)+in1(i)*in2(i)
  end do
!$OMP END DO
!$OMP END PARALLEL
end function
```

```
double dot_product(int n, float *in1, *in2)
{
  int i,ith;
  double out=0.;
  #pragma omp parallel default(shared) private(ith,i)
    reduction(+:out)
    ith=omp_get_thread_num();
  #pragma omp for
  {
    for(i=0; i < n; i++){
      out+=in1[i]*in2[i];
    }
  }
  return(out);
}
```

# Dot product (version 3)

```
double precision function dot_product(in1,in2)
  real :: in1(:),in2(:)
  integer :: i,ith
  integer,external :: omp_get_thread_num
  dots=0;dot_product=0
  !$OMP PARALLEL default(shared) private(i)
  !$OMP REDUCTION(+:dot_product)
  ith=omp_get_thread_num()
  !$OMP DO
  do i=1,size(in1)
    dot_product=dots(ith+1)+in1(i)*in2(i)
  end do
  !$OMP END DO
  !$OMP END PARALLEL
end function
```

```
double dot_product(int n, float *in1, *in2)
{
  int i,ith;
  double out=0.;
  #pragma omp parallel default(shared) private(ith,i)
  reduction(+:out)
  ith=omp_get_thread_num();
  #pragma omp for
  {
    for(i=0; i < n; i++){
      out+=in1[i]*in2[i];
    }
  }
  return(out);
}
```

All values are combined at the end of the parallel clause.

# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i)
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i)
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

Test: Four bugs/differences in this conversion



# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i)
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

Problem I:  
a and b are clobbered within  
the do loop

# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i,a,b)
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

**Solution 1:**  
Make a and b private variables

# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i,a,b)
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

**Problem 2:**  
a is initialized to 0 when entering parallel  
environment

# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i,a,b)&
!$OMP & firstprivate(a)
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

Solution 2:  
firstprivate clause copies variable from  
serial portion to parallel region

# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i,a,b)&
!$OMP & firstprivate(a)
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

Problem 3: b disappears when leaving the parallel environment.

# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i,a,b)&
!$OMP & firstprivate(a) lastprivate(b)
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

Solution: lastprivate copies from the last parallel thread back to the serial code portion.

# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i,a,b)&
!$OMP & firstprivate(a) lastprivate(b)
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

Problem: Print statements will not be in  
the same order

# Serial to OpenMP gotchas

```
a=4
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  print i,a,b
  call big_stuff2(array,i)
end do
print b
```

```
a=4
!$OMP PARALLEL DO default(shared) private(i,a,b)&
!$OMP & firstprivate(a) lastprivate(b) ordered
do i=1,n
  call big_stuff(array,i)
  a=a-1
  b=a+i
  !$OMP ORDERED
  print i,a,b
  !$OMP END ORDERED
  call big_stuff2(array,i)
end do
!$OMP END PARALLEL DO
print b
```

Solution: Ordered  
Guarantees that code block will be  
executed in the same manner as serial  
code



# OMP nowait

```
for(i=0; i < n; i++)  
    c[i]=a[i]*a[i];
```

```
for(i=0; i < n; i++)  
    d[i]=b[i]*b[i];
```

```
#pragma omp parallel  
{  
    #pragma omp for private(i)  
    for(i=0; i < n; i++)  
        c[i]=a[i]*a[i];  
  
    #pragma omp for private(i)  
    for(i=0; i < n; i++)  
        d[i]=b[i]*b[i];  
}
```

End of the parallel block implies a barrier statement.

# OMP nowait

```
for(i=0; i < n; i++)  
    c[i]=a[i]*a[i];
```

```
for(i=0; i < n; i++)  
    d[i]=b[i]*b[i];
```

```
#pragma omp parallel nowait  
{  
    #pragma omp for private(i)  
    for(i=0; i < n; i++)  
        c[i]=a[i]*a[i];  
  
    #pragma omp for private(i)  
    for(i=0; i < n; i++)  
        d[i]=b[i]*b[i];  
}
```

If loops aren't dependent the nowait clause removes the barrier statement.

# OMP workshare

```
real a(:),b(:),c(:)
```

```
c=b+a
```

```
real a(:),b(:),c(:)
```

```
#OMP PARALLEL WORKSHARE
```

```
c=b+a
```

```
#OMP END PARALLEL WORKSHARE
```

Parallelize Fortran array operations.

# Threads model

program alpha

.

.

call sub1()

call sub2()

call sub3()

call sub4()

call sub5()

.

.

end program

Time



# Threads model

program alpha

.

.

call sub1()

call sub2()

call sub3()

call sub4()

call sub5()

.

.

end program

Thread 1



Thread 2



Thread 3



Thread 4



Thread 5



Time



# Sections

```
program alpha  
.  
.  
!$OMP PARALLEL SECTIONS  
call sub1()  
call sub2()  
call sub3()  
call sub4()  
call sub5()  
!$END PARALLEL SECTIONS  
.  
end program
```

Begin and end functional parallel  
environment


# Sections

```
program alpha
.  
.  
!$OMP PARALLEL SECTIONS  
!$OMP SECTION  
call sub1()  
!$OMP END SECTION  
!$OMP SECTION  
call sub2()  
!$OMP END SECTION  
!$OMP SECTION  
call sub3()  
!$OMP END SECTION  
!$OMP SECTION  
call sub4()  
!$OMP END SECTION  
!$OMP SECTION  
call sub5()  
!$OMP END SECTION  
!$END PARALLEL SECTIONS  
.  
end program
```

Only a single thread will enter each  
section clause

# Single processor region/1

*This construct is ideally suited for I/O or initialization*



```
for (i=0; i < N; i++)  
{  
    .....  
    "read a[0..N-1]";  
    .....  
}
```

*Serial*

*"declare A to be shared"*

```
#pragma omp parallel for  
for (i=0; i < N; i++)  
{  
    .....  
    .....  
    one volunteer requested  
    .....  
    "read a[0..N-1]";  
    .....  
    thanks, we're done  
    .....  
}
```



May have to insert a  
barrier here

*Parallel*



# Single processor region/2

- *Usually, there is a barrier needed after this region*
- *Might therefore be a scalability bottleneck (Amdahl's law)*



# SINGLE and MASTER construct

*Only one thread in the team executes the code enclosed*

```
#pragma omp single [clause[[,] clause] ...]  
{  
    <code-block>  
}
```

```
!$omp single [clause[[,] clause] ...]  
    <code-block>  
!$omp end single [nowait]
```

*Only the master thread executes the code block:*

```
#pragma omp master  
{ <code-block> }
```

```
!$omp master  
    <code-block>  
!$omp end master
```

*There is no implied  
barrier on entry or  
exit !*

# More synchronization directives

*The enclosed block of code is executed in the order in which iterations would be executed sequentially:*

```
#pragma omp ordered  
{ <code-block> }
```

```
!$omp ordered  
    <code-block>  
!$omp end ordered
```

**Expensive !**

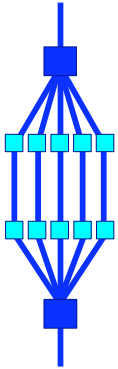
*Ensure that all threads in a team have a consistent view of certain objects in memory:*

```
#pragma omp flush [(list)]
```

```
!$omp flush [(list)]
```

**In the absence of a list,  
all visible variables are  
flushed**

# OpenMP environment variables




OpenMP environment variable	Default for Sun OpenMP
OMP_NUM_THREADS <u>n</u>	1
OMP_SCHEDULE “ <u>schedule</u> ,[ <u>chunk</u> ]”	static, “N/P” (1)
OMP_DYNAMIC { TRUE   FALSE }	TRUE (2)
OMP_NESTED { TRUE   FALSE }	FALSE (3)

- (1) *The chunk size approximately equals the number of iterations (N) divided by the number of threads (P)*
- (2) *The number of threads will be limited to the number of on-line processors in the system. This can be changed by setting **OMP\_DYNAMIC** to **FALSE**.*
- (3) *Multi-threaded execution of inner parallel regions in nested parallel regions is supported as of Sun Studio 10*

**Note: The names are in uppercase, the values are case insensitive**

# Global data - example



```

      ....
      include "global.h"
      ....
!$omp parallel private(j)
  do j = 1, n
    call suba(j)
  end do
!$omp end do
!$omp end parallel
      ....

```

*file global.h*

```

common /work/a(m,n),b(m)

```

```

subroutine suba(j)
  ....
  include "global.h"
  ....

```

```

do i = 1, m
  b(i) = j
end do

```

**Race condition !**


```

do i = 1, m
  a(i,j) = func_call(b(i))
end do

return
end

```

# Global data - race condition



*Thread 1*



call suba(1)

*Thread 2*



call suba(2)

Shared

subroutine suba(j=1)

```
do i = 1, m
  b(i) = 1
end do
```


```
.....
do i = 1, m
  a(i,1)=func_call(b(i))
end do
```

subroutine suba(j=2)

```
do i = 1, m
  b(i) = 2
end do
```

```
.....
do i = 1, m
  a(i,2)=func_call(b(i))
end do
```

# Example - solution



```

      ....
      include "global.h"
      ....
!$omp parallel private(j)
      do j = 1, n
          call suba(j)
      end do
!$omp end do
!$omp end parallel
      ....

```

*new file global.h*

```

integer, parameter:: nthreads=4
common /work/a(m,n)
common /tprivate/b(m,nthreads)

```

```

subroutine suba(j)
      ....
      include "global.h"
      ....

      TID = omp_get_thread_num()+1
      do i = 1, m
          b(i,TID) = j
      end do


      do i = 1, m
          a(i,j)=func_call(b(i,TID))
      end do

      return
end

```

- ☞ *By expanding array B, we can give each thread unique access to it's storage area*
- ☞ *Note that this can also be done using dynamic memory (allocatable, malloc, ....)*

# Example - solution 2



```

      ....
      include "global.h"
      ....
!$omp parallel private(j)
      do j = 1, n
          call suba(j)
      end do
!$omp end do
!$omp end parallel
      ....

```

*new file global.h*

```

common /work/a(m,n)
common /tprivate/b(m)
!$omp threadprivate(/tprivate/)

```

```

subroutine suba(j)
      ....
      include "global.h"
      ....

      do i = 1, m
          b(i) = j
      end do

      do i = 1, m
          a(i,j) = func_call(b(i))
      end do


      return
end

```

- ☞ *The compiler will create thread private copies of array B, to give each thread unique access to it's storage area*
- ☞ *Note that the number of copies will be automatically adapted to the number of threads*



# About global data



- ❑ Global data is shared and requires special care
- ❑ *A problem may arise in case multiple threads access the same memory section simultaneously:*
  - *Read-only data is no problem*
  - *Updates have to be checked for race conditions*
- ❑ *It is your responsibility to deal with this situation*
- ❑ *In general one can do the following:*
  - *Split the global data into a part that is accessed in serial parts only and a part that is accessed in parallel*
  - *Manually create thread private copies of the latter*
  - *Use the thread ID to access these private copies*
- ❑ **Alternative: Use OpenMP's threadprivate construct**

# The threadprivate construct

## □ *OpenMP's threadprivate directive*

```
!$omp threadprivate (/cb/ [,/cb/] ...)
```


```
#pragma omp threadprivate (list)
```

## □ *Thread private copies of the designated global variables and common blocks will be made*

## □ *Several restrictions and rules apply when doing this:*

- *The number of threads has to remain the same for all the parallel regions (i.e. no dynamic threads)*
  - ✓ *Sun implementation supports changing the number of threads*
- *Initial data is undefined, unless **copyin** is used*
- *.....*

## □ *Check the documentation when using threadprivate !*



# The copyin clause

## copyin (list)

- ✓ *Applies to THREADPRIVATE common blocks only*
- ✓ *At the start of the parallel region, data of the master thread is copied to the thread private copies*

### Example:


```
common /cblock/velocity
common /fields/xfield, yfield, zfield

! create thread private common blocks

!$omp threadprivate (/cblock/, /fields/)

!$omp parallel          &
!$omp default (private) &
!$omp copyin ( /cblock/, zfield )
```

# Runtime library overview




## *Name*

*omp\_set\_num\_threads*  
*omp\_get\_num\_threads*  
*omp\_get\_max\_threads*  
*omp\_get\_thread\_num*  
*omp\_get\_num\_procs*  
*omp\_in\_parallel*  
*omp\_set\_dynamic*  
  
*omp\_get\_dynamic*  
*omp\_set\_nested*  
  
*omp\_get\_nested*  
*omp\_get\_wtime*  
*omp\_get\_wtick*

## *Functionality*


*Set number of threads*  
*Return number of threads in team*  
*Return maximum number of threads*  
*Get thread ID*  
*Return maximum number of processors*  
*Check whether in parallel region*  
*Activate dynamic thread adjustment*  
*(but implementation is free to ignore this)*  
*Check for dynamic thread adjustment*  
*Activate nested parallelism*  
*(but implementation is free ignore this)*  
*Check for nested parallelism*  
*Returns wall clock time*  
*Number of seconds between clock ticks*

# OpenMP locking routines



- ❑ *Locks provide greater flexibility over critical sections and atomic updates:*
  - *Possible to implement asynchronous behaviour*
  - *Not block structured*
- ❑ *The so-called lock variable, is a special variable:*
  - *Fortran: type INTEGER and of a KIND large enough to hold an address*
  - *C/C++: type `omp_lock_t` and `omp_nest_lock_t` for nested locks*
- ❑ *Lock variables should be manipulated through the API only*
- ❑ *It is illegal, and behaviour is undefined, in case a lock variable is used without the appropriate initialization*

# Nested locking



- ❑ *Simple locks: may not be locked if already in a locked state*
- ❑ *Nestable locks: may be locked multiple times by the same thread before being unlocked*
- ❑ *In the remainder, we will discuss simple locks only*
- ❑ *The interface for functions dealing with nested locks is similar (but using nestable lock variables):*

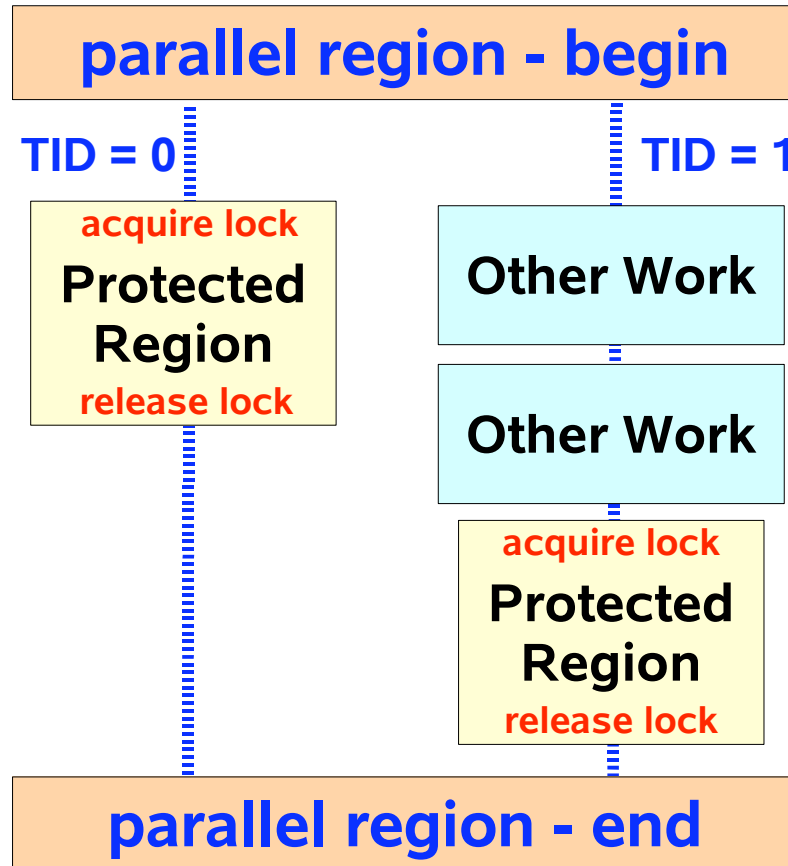

## Simple locks

```
omp_init_lock  
omp_destroy_lock  
omp_set_lock  
omp_unset_lock  
omp_test_lock
```

## Nestable locks


```
omp_init_nest_lock  
omp_destroy_nest_lock  
omp_set_nest_lock  
omp_unset_nest_lock  
omp_test_nest_lock
```

# OpenMP locking example



- ♦ *The protected region contains the update of a shared variable*
- ♦ *One thread will acquire the lock and perform the update*
- ♦ *Meanwhile, the other thread will do some other work*
- ♦ *When the lock is released again, the other thread will perform the update*

# Locking example - the code



Program Locks

....

Call `omp_init_lock (LCK)`

Initialize lock variable

`!$omp parallel shared(SUM,LCK) private(TID)`

`TID = omp_get_thread_num()`

Check availability of lock  
(will also set the lock)

Do While ( `omp_test_lock (LCK)` .EQV. .FALSE. )

*Call Do\_Something\_Else(TID)*

End Do

*Call Do\_Work(SUM,TID)*

Release lock again

Call `omp_unset_lock (LCK)`

`!$omp end parallel`

Remove lock association


Call `omp_destroy_lock (LCK)`

Stop

End



# Example output for 2 threads



```

TID: 1 at 09:07:27 => entered parallel region
TID: 1 at 09:07:27 => done with WAIT loop and has the lock
TID: 1 at 09:07:27 => ready to do the parallel work
TID: 1 at 09:07:27 => this will take about 18 seconds
TID: 0 at 09:07:27 => entered parallel region
TID: 0 at 09:07:27 => WAIT for lock - will do something else for 5 seconds
TID: 0 at 09:07:32 => WAIT for lock - will do something else for 5 seconds
TID: 0 at 09:07:37 => WAIT for lock - will do something else for 5 seconds
TID: 0 at 09:07:42 => WAIT for lock - will do something else for 5 seconds
TID: 1 at 09:07:45 => done with my work
TID: 1 at 09:07:45 => done with work loop - released the lock
TID: 1 at 09:07:45 => ready to leave the parallel region
TID: 0 at 09:07:47 => done with WAIT loop and has the lock
TID: 0 at 09:07:47 => ready to do the parallel work
TID: 0 at 09:07:47 => this will take about 18 seconds
TID: 0 at 09:08:05 => done with my work
TID: 0 at 09:08:05 => done with work loop - released the lock
TID: 0 at 09:08:05 => ready to leave the parallel region
Done at 09:08:05 - value of SUM is 1100

```

Used to check the answer

*Note: program has been instrumented to get this information*

# Bug examples

```
int i, chunk, tid;
float a[N], b[N], c[N];

/* Some initializations */
for (i=0; i < N; i++)
    a[i] = b[i] = i * 1.0;;

#pragma omp parallel for    \
    shared(a,b,c,chunk)    \
    private(i,tid)
{
    tid = omp_get_thread_num();
    for (i=0; i < N; i++)
    {
        c[i] = a[i] + b[i];
        printf("tid= %d i= %d c[i]= %f\n", tid, i, c[i]);
    }
} /* end of parallel for construct */

}
```

# Bug examples

```
int i, chunk, tid;
float a[N], b[N], c[N];

/* Some initializations */
for (i=0; i < N; i++)
    a[i] = b[i] = i * 1.0;;

#pragma omp parallel for    \
    shared(a,b,c,chunk)    \
    private(i,tid)
{
    tid = omp_get_thread_num();
    for (i=0; i < N; i++)
    {
        c[i] = a[i] + b[i];
        printf("tid= %d i= %d c[i]= %f\n", tid, i, c[i]);
    }
} /* end of parallel for construct */

}
```

# Bug examples

```
int i, chunk, tid;
float a[N], b[N], c[N];

/* Some initializations */
for (i=0; i < N; i++)
    a[i] = b[i] = i * 1.0;;

#pragma omp parallel for \
    shared(a,b,c,chunk) \
    private(i,tid)
{
    tid = omp_get_thread_num();
    for (i=0; i < N; i++)
    {
        c[i] = a[i] + b[i];
        printf("tid= %d i= %d c[i]= %f\n", tid, i, c[i]);
    }
} /* end of parallel for construct */

}
```

```
int i, chunk, tid;
float a[N], b[N], c[N];
char first_time;

/* Some initializations */
for (i=0; i < N; i++)
    a[i] = b[i] = i * 1.0;
first_time = 'y';

#pragma omp parallel for \
    shared(a,b,c,chunk) \
    private(i,tid) \
    firstprivate(first_time)

for (i=0; i < N; i++)
{
    if (first_time == 'y')
    {
        tid = omp_get_thread_num();
        first_time = 'n';
    }
    c[i] = a[i] + b[i];
    printf("tid= %d i= %d c[i]= %f\n", tid, i, c[i]);
}

}
```

# Bug examples

```
int nthreads, i, tid;
float total;

/** Spawn parallel region **/
#pragma omp parallel
{
    /* Obtain thread number */
    tid = omp_get_thread_num();
    /* Only master thread does this */
    if (tid == 0) {
        nthreads = omp_get_num_threads();
        printf("Number of threads = %d\n", nthreads);
    }
    printf("Thread %d is starting...\n",tid);

    #pragma omp barrier

    /* do some work */
    total = 0.0;
    #pragma omp for
    for (i=0; i<1000000; i++)
        total = total + i*1.0;

    printf ("Thread %d is done! Total= %e\n",tid,total);
} /** End of parallel region **/
}
```

<http://www.llnl.gov/computing/tutorials/openMP/samples>

# Bug examples

```
int nthreads, i, tid;
float total;

/** Spawn parallel region **/
#pragma omp parallel
{
    /* Obtain thread number */
    tid = omp_get_thread_num();
    /* Only master thread does this */
    if (tid == 0) {
        nthreads = omp_get_num_threads();
        printf("Number of threads = %d\n", nthreads);
    }
    printf("Thread %d is starting...\n",tid);

    #pragma omp barrier

    /* do some work */
    total = 0.0;
    #pragma omp for
    for (i=0; i<1000000; i++)
        total = total + i*1.0;

    printf ("Thread %d is done! Total= %e\n",tid,total);

} /** End of parallel region **/
}
```

<http://www.llnl.gov/computing/tutorials/openMP/samples>

# Bug examples

```
int nthreads, i, tid;
float total;

/** Spawn parallel region ***/
#pragma omp parallel
{
    /* Obtain thread number */
    tid = omp_get_thread_num();
    /* Only master thread does this */
    if (tid == 0) {
        nthreads = omp_get_num_threads();
        printf("Number of threads = %d\n", nthreads);
    }
    printf("Thread %d is starting...\n",tid);

    #pragma omp barrier

    /* do some work */
    total = 0.0;
    #pragma omp for
    for (i=0; i<1000000; i++)
        total = total + i*1.0;

    printf ("Thread %d is done! Total= %e\n",tid,total);
} /** End of parallel region ***/
}
```

<http://www.llnl.gov/computing/tutorials/openMP/samples>

```
int nthreads, i, tid;
float total;

/** Spawn parallel region ***/
#pragma omp parallel
{
    /* Obtain thread number */
    tid = omp_get_thread_num();
    /* Only master thread does this */
    if (tid == 0) {
        nthreads = omp_get_num_threads();
        printf("Number of threads = %d\n", nthreads);
    }
    printf("Thread %d is starting...\n",tid);

    #pragma omp barrier

    /* do some work */
    total = 0.0;
    #pragma omp for reduction(+:total)
    for (i=0; i<1000000; i++)
        total = total + i*1.0;

    printf ("Thread %d is done! Total= %e\n",tid,total);

} /** End of parallel region ***/
}
```

# Segmentation fault

```
#define N 1048
int main (int argc, char *argv[]) {
    int nthreads, tid, i, j;
    double a[N][N];

    /* Fork a team of threads with explicit variable scoping */
    #pragma omp parallel shared(nthreads) private(i,j,tid,a)
    {

        /* Obtain/print thread info */
        tid = omp_get_thread_num();
        if (tid == 0)
        {
            nthreads = omp_get_num_threads();
            printf("Number of threads = %d\n", nthreads);
        }
        printf("Thread %d starting...\n", tid);

        /* Each thread works on its own private copy of the array */
        for (i=0; i<N; i++)
            for (j=0; j<N; j++)
                a[i][j] = tid + i + j;
        /* For confirmation */
        printf("Thread %d done. Last element= %f\n",tid,a[N-1][N-1]);
    } /* All threads join master thread and disband */
}
```



# Bug examples

```
#define N 1048
int main (int argc, char *argv[]) {
int nthreads, tid, i, j;
double a[N][N];
```

```
/* Fork a team of threads with explicit variable scoping */
#pragma omp parallel shared(nthreads) private(i,j,tid,a)
{
```

```
/* Obtain/print thread info */
```

```
tid = omp_get_thread_num();
```

```
if (tid == 0)
```

```
{
nthreads = omp_get_num_threads();
printf("Number of threads = %d\n", nthreads);
}
```

```
printf("Thread %d starting...\n", tid);
```

```
/* Each thread works on its own private copy of the array */
```

```
for (i=0; i<N; i++)
```

```
for (j=0; j<N; j++)
```

```
a[i][j] = tid + i + j;
```

```
/* For confirmation */
```

```
printf("Thread %d done. Last element= %f\n",tid,a[N-1][N-1]);
```

```
} /* All threads join master thread and disband */
```

```
}
```

```
#define N 1048
int main (int argc, char *argv[]) {
int nthreads, tid, i, j;
double ***a;
```

```
#pragma omp parallel private(nthreads)
```

```
nthreads=omp_get_num_threads();
```

```
a=(double***) allocate3(n,n,nth);
```

```
/* Fork a team of threads with explicit variable scoping */
```

```
#pragma omp parallel shared(nthreads,a) private(i,j,tid)
```

```
{
```

```
/* Obtain/print thread info */
```

```
tid = omp_get_thread_num();
```

```
printf("Thread %d starting...\n", tid);
```

```
/* Each thread works on its own private copy of the array */
```

```
for (i=0; i<N; i++)
```

```
for (j=0; j<N; j++)
```

```
a[tid][i][j] = tid + i + j;
```

```
/* For confirmation */
```

```
printf("Thread %d done. Last element= %f\n",tid,a[N-1][N-1]);
```

```
} /* All threads join master thread and disband */
```

```
}
```

# Serial code

```
float alpha;
void main(int argc, char **argv){
.
.
.
in=(float**)alloc_2d_float(n1,n2);
work=(float*)alloc_1d_float(n1);
out=(float**) alloc_2d_float(n1,n2);

for(i=0; i < n2; i++){
    for(j=0; j < n1; j++){
        .
        .
        alpha=func(in[i][j]....)
        work[j]=func(in[i][j].....)
        .
        .
        out[i][j]=func(work[j]...)
    }
}
```

# Determine where to parallelize

```
float alpha;
void main(int argc, char **argv){
.
.
.
in=(float**)alloc_2d_float(n1,n2);
work=(float*)alloc_1d_float(n1);
out=(float**) alloc_2d_float(n1,n2);

for(i=0; i < n2; i++){
    for(j=0; j < n1; j++){
        .
        .
        alpha=func(in[i][j]....)
        work[j]=func(in[i][j]....)
        .
        .
        out[i][j]=func(work[j]...)
    }
}
```

# Put serial portion in separate function

```
float alpha;
void main(int argc, char **argv){
.
.
.
in=(float**)alloc_2d_float(n1,n2);
work=(float*)alloc_1d_float(n1);
out=(float**) alloc_2d_float(n1,n2);

for(i=0; i < n2; i++){
    serial_work(in,out,work,i);
}
```

```
void serial_work(float **in, float **out, float *work, int i){
.
    for(j=0; j < n1; j++){
        .
        .
        alpha=func(in[i][j]....)
        work[j]=func(in[i][j].....)
        .
        .
        out[i][j]=func(work[j]...)
    }
}
```

# Get rid of global variables modified in serial portion

```
float alpha;
void main(int argc, char **argv){
.
.
.
in=(float**)alloc_2d_float(n1,n2);
work=(float*)alloc_1d_float(n1);
out=(float**) alloc_2d_float(n1,n2);

for(i=0; i < n2; i++){
    serial_work(in,out,work,i);
}
```

```
void serial_work(float **in, float **out, float *work, int i){
.
for(j=0; j < n1; j++){
.
.
    alpha=func(in[i][j]....)
    work[j]=func(in[i][j].....)
.
.
    out[i][j]=func(work[j]...)
}
}
```

# Get rid of global variables modified in serial portion

```
void main(int argc, char **argv){  
.  
.  
.  
in=(float**)alloc_2d_float(n1,n2);  
work=(float*)alloc_1d_float(n1);  
out=(float**) alloc_2d_float(n1,n2);  
  
for(i=0; i < n2; i++){  
    serial_work(in,out,work,i);  
}
```

```
void serial_work(float **in, float **out, float *work, int i){  
    float alpha;  
    .  
    for(j=0; j < n1; j++){  
        .  
        .  
        alpha=func(in[i][j]....)  
        work[j]=func(in[i][j].....)  
        .  
        .  
        out[i][j]=func(work[j]...)  
    }  
}
```

# Find intermediate heap arrays

```
void main(int argc, char **argv){
.
.
.
in=(float**)alloc_2d_float(n1,n2);
work=(float*)alloc_1d_float(n1);
out=(float**) alloc_2d_float(n1,n2);

for(i=0; i < n2; i++){
    serial_work(in,out,work,i);
}
```

```
void serial_work(float **in, float **out, float *work, int i){
    float alpha;
.
    for(j=0; j < n1; j++){
        .
        .
        alpha=func(in[i][j]....)
        work[j]=func(in[i][j].....)
        .
        .
        out[i][j]=func(work[j]...)
    }
}
```

# Expand dimensionality by nthreads

```
void main(int argc, char **argv){
```

```
·  
·
```

```
    nth=1; //For serial case
```

```
    in=(float**)alloc_2d_float(n1,n2);
```

```
    work=(float**)alloc_2d_float(n1,nth);
```

```
    out=(float**) alloc_2d_float(n1,n2);
```

```
    ith=0;
```

```
    for(i=0; i < n2; i++){
```

```
        serial_work(in,out,work,i,ith);
```

```
    }
```

```
void serial_work(float **in, float **out, float **work, int i,  
int ith){  
    float alpha;
```

```
·  
    for(j=0; j < n1; j++){
```

```
·
```

```
·
```

```
    alpha=func(in[i][j]....)
```

```
    work[ith][j]=func(in[i][j].....)
```

```
·
```

```
·
```

```
    out[i][j]=func(work[ith][j]...)
```

```
    }
```

```
}
```



# Figure out number of threads

```
void main(int argc, char **argv){  
.  
.  
  
nth=1; //For serial case  
#pragma omp parallel  
nth=omp_get_num_threads();  
#pragma omp end parallel  
  
in=(float**)alloc_2d_float(n1,n2);  
work=(float**)alloc_2d_float(n1,nth);  
out=(float**) alloc_2d_float(n1,n2);  
  
ith=0;  
for(i=0; i < n2; i++){  
    serial_work(in,out,work,i,ith);  
}
```

```
void serial_work(float **in, float **out, float **work, int i,  
int ith){  
    float alpha;  
    .  
    for(j=0; j < n1; j++){  
        .  
        .  
        alpha=func(in[i][j]....)  
        work[ith][j]=func(in[i][j].....)  
        .  
        .  
        out[i][j]=func(work[ith][j]...)  
    }  
}
```

# Parallelize loop

```
void main(int argc, char **argv){
.
.

nth=1; //For serial case
#pragma omp parallel
nth=omp_get_num_threads();
#pragma omp end parallel

in=(float**)alloc_2d_float(n1,n2);
work=(float**)alloc_2d_float(n1,nth);
out=(float**) alloc_2d_float(n1,n2);

ith=0;
#pragma omp parallel for private(i)
for(i=0; i < n2; i++){
    serial_work(in,out,work,i,ith);
}
```

```
void serial_work(float **in, float **out, float **work, int i,
int ith){
float alpha;
.
.
for(j=0; j < n1; j++){
.
.
alpha=func(in[i][j]....)
work[ith][j]=func(in[i][j].....)
.
.
out[i][j]=func(work[ith][j]...)
}
}
```

# Grab thread number


```
void main(int argc, char **argv){  
.  
.  
  
nth=1; //For serial case  
#pragma omp parallel  
nth=omp_get_num_threads();  
#pragma omp end parallel  
  
in=(float**)alloc_2d_float(n1,n2);  
work=(float**)alloc_2d_float(n1,nth);  
out=(float**) alloc_2d_float(n1,n2);  
  
#pragma omp parallel for private(i)  
for(i=0; i < n2; i++){  
    ith=omp_get_thread_num();  
    serial_work(in,out,work,i,ith);  
}
```

```
void serial_work(float **in, float **out, float **work, int i,  
int ith){  
    float alpha;  
.  
    for(j=0; j < n1; j++){  
        .  
        .  
        alpha=func(in[i][j]....)  
        work[ith][j]=func(in[i][j].....)  
        .  
        .  
        out[i][j]=func(work[ith][j]...)  
    }  
}
```


# More complex examples

- Matrix-vector multiply
- Sparse matrix vector multiply
- Matrix-matrix multiply
- Convolution

# Matrix-vector multiply




# Matrix-vector multiply



# Serial code

```
do i2=1,n2
  do i1=1,n1
    rsum=rsum+buf(i1,i2)*vecin(i1)
  end do
  vecout(i2)=rsum
end do
```

# Parallelize over inner axis







# Parallelize over inner access

```
do i2=1,n2
  !$OMP PARALLEL DO private(i1) reduction (+:rsum)
  do i1=1,n1
    rsum=rsum+buf(i1,i2)*vecin(i1)
  end do
  !OMP END PARALLEL DO
  vecout(i2)=rsum
end do
```

# Speedup




# Parallelize over outer axis




# Parallelize over inner access

```
!$OMP PARALLEL DO private(rsum,i2,il)
do i2=1,n2
  rsum=0
  do il=1,nl
    rsum=rsum+buf(il,i2)*vecin(il)
  end do
  vecout(i2)=rsum
end do
!OMP END PARALLEL DO
```

# Speedup



# Parallelize over outer axis



# Blocking


```
i=0
do i2=1,nbig2
  do i1=1,nbig1
    i=i+1
    ibeg(:,i)=(/nbuf1*(i1-1),nbuf2*(i2-1)/)
  end do
end do

!$OMP PARALLEL DO private(i1) schedule(dynamic)
  do i=1,nbig1*nbig2
    call sub_mult(buf,vecin,vecout,nbuf1,nbuf2,ibeg(1,i),ibeg(2,i))
  end do
!OMP END PARALLEL DO
```

```
subroutine sub_mult(buf,vecin,vecout,n1,n2,i1beg,i2beg)
  real :: buf(:,,:),vecin(:),vecout(:)
  integer :: n1,n2,i1beg,i2beg
  integer :: i1,i2
  real :: rsum


  do i2=1,n2
    rsum=0
    do i1=1,n1
      rsum=rsum+buf(i1+i1beg,i2+i2beg)
    end do
    !OMP ATOMIC
    vecout(i2)=vecout(i2)+rsum
  end do
end subroutine
```

# Speedup







# Dynamic blocking




# Speedup



# Buffer size




# Matrix-vector multiply




out

=




buf1

\*




buf2

# Matrix-matrix multiply




out

=



buf1

\*




buf2

# Matrix multiplication


```
do i3=1,n
  do i2=1,n
    rsum=0
    do i1=1,n
      rsum=rsum+in1(i1,i3)*in2(i2,i1)
    end do
    out(i2,i3)=rsum
  end do
end do
```

# Matrix-vector multiply




out

=



buf1

\*




buf2

# Matrix multiplication

```
do i3=1,n
  do i2=1,n
    rsum=0
    !OMP PARALLEL DO private(i1) reduction(+:rsum)
    do i1=1,n
      rsum=rsum+in1(i1,i3)*in2(i2,i1)
    end do
    !OMP END PARALLEL DO
    out(i2,i3)=rsum
  end do
end do
```




# Matrix-vector multiply




out

=



buf1

\*




buf2

# Matrix multiplication


```
do i3=1,n
  !OMP PARALLEL DO private(rsum,i2,i1)
  do i2=1,n
    rsum=0
    do i1=1,n
      rsum=rsum+in1(i1,i3)*in2(i2,i1)
    end do
    out(i2,i3)=rsum
  end do
  !OMP END PARALLEL DO
end do
```

# Matrix-vector multiply




out

=



buf1

\*




buf2


# Matrix multiplication

```
!$OMP PARALLEL DO private(rsum,i1,i2,i3)
do i3=1,n
  do i2=1,n
    rsum=0
    do i1=1,n
      rsum=rsum+in1(i1,i3)*in2(i2,i1)
    end do
    out(i2,i3)=rsum
  end do
end do
!OMP END PARALLEL DO
```


# Matrix multiply




# Matrix multiply



# Matrix multiply




# Matrix multiply






# Matrix multiply




# Buffer size




# Convolution

```
for(it=0; it < nt; it++){  
    do_convolution  
    pm=p; p=pp  
}
```

```
for (ix=1; ix<nx-1; ++ix) {  
    for (iz=1; iz<nz-1; ++iz) {  
        pp[ix][iz] = 2.0*p[ix][iz]-pm[ix][iz] +  
            scale*dvv[ix][iz]*( p[ix+1][iz]+ p[ix-1][iz]+  
            p[ix][iz+1]+ p[ix][iz-1]-4.0*p[ix][iz])+s[ix][iz]  
    }  
}
```



# Convolution: Larger stencil and 3-D




# Convolution

```
for(iy=3;iy < ny-4; ++iy){
  for (ix=3; ix<nx-4; ++ix) {
    for (iz=3; iz<nz-4; ++iz) {
      pp[iy][ix][iz] = 2.0*p[iy][ix][iz]-pm[iy][ix][iz] +
        scale*dvv[ix][iz]*(
          a3*p[iy-3][ix][iz]+a2*p[iy-2][ix][iz]+a1*p[iy-1][ix][iz]+a0*p[iy][ix][iz]+
          a1*p[iy+1][ix][iz]+a2*p[iy+2][ix][iz]+a3*p[iy+3][ix][iz]+
          a3*p[iy][ix-3][iz]+a2*p[iy][ix-2][iz]+a1*p[iy][ix-1][iz]+a0*p[iy][ix][iz]+
          a1*p[iy][ix+1][iz]+a2*p[iy][ix+2][iz]+a3*p[iy][ix+3][iz]+
          a3*p[iy][ix][iz-3]+a2*p[iy][ix][iz-2]+a1*p[iy][ix][iz-1]+a0*p[iy][ix][iz]+
          a1*p[iy][ix][iz+1]+a2*p[iy][ix][iz+2]+a3*p[iy][ix][iz+3]));
    }
  }
}
```

# Convolution

```
for(iy=3;iy < ny-4; ++iy){
for (ix=3; ix<nx-4; ++ix) {
    for (iz=3; iz<nz-4; ++iz) {
        pp[iy][ix][iz] = 2.0*p[iy][ix][iz]-pm[iy][ix][iz] +
            scale*dvw[ix][iz]*(
                a3*p[iy-3][ix][iz]+a2*p[iy-2][ix][iz]+a1*p[iy-1][ix][iz]+a0*p[iy][ix][iz]+
                a1*p[iy+1][ix][iz]+a2*p[iy+2][ix][iz]+a3*p[iy+3][ix][iz]+
                a3*p[iy][ix-3][iz]+a2*p[iy][ix-2][iz]+a1*p[iy][ix-1][iz]+a0*p[iy][ix][iz]+
                a1*p[iy][ix+1][iz]+a2*p[iy][ix+2][iz]+a3*p[iy][ix+3][iz]+
                a3*p[iy][ix][iz-3]+a2*p[iy][ix][iz-2]+a1*p[iy][ix][iz-1]+a0*p[iy][ix][iz]+
                a1*p[iy][ix][iz+1]+a2*p[iy][ix][iz+2]+a3*p[iy][ix][iz+3]));
    }
}
}
```

# Convolution




# Convolution

```
for(iy=3;iy < ny-4; ++iy){  
  for (ix=3; ix<nx-4; ++ix) {  
    for (iz=3; iz<nz-4; ++iz) {  
      pp[iy][ix][iz] = 2.0*p[iy][ix][iz]-pm[iy][ix][iz] +  
        scale*dvw[ix][iz]*(  
          a3*p[iy-3][ix][iz]+a2*p[iy-2][ix][iz]+a1*p[iy-1][ix][iz]+a0*p[iy][ix][iz]+  
          a1*p[iy+1][ix][iz]+a2*p[iy+2][ix][iz]+a3*p[iy+3][ix][iz]+  
          a3*p[iy][ix-3][iz]+a2*p[iy][ix-2][iz]+a1*p[iy][ix-1][iz]+a0*p[iy][ix][iz]+  
          a1*p[iy][ix+1][iz]+a2*p[iy][ix+2][iz]+a3*p[iy][ix+3][iz]+  
          a3*p[iy][ix][iz-3]+a2*p[iy][ix][iz-2]+a1*p[iy][ix][iz-1]+a0*p[iy][ix][iz]+  
          a1*p[iy][ix][iz+1]+a2*p[iy][ix][iz+2]+a3*p[iy][ix][iz+3]));  
    }  
  }  
}
```




# Convolution



# Convolution

```
for(iy=3;iy < ny-4; ++iy){
  for (ix=3; ix<nx-4; ++ix) {
    for (iz=3; iz<nz-4; ++iz) {
      pp[iy][ix][iz] = 2.0*p[iy][ix][iz]-pm[iy][ix][iz] +
        scale*dvw[ix][iz]*(
          a3*p[iy-3][ix][iz]+a2*p[iy-2][ix][iz]+a1*p[iy-1][ix][iz]+a0*p[iy][ix][iz]+
          a1*p[iy+1][ix][iz]+a2*p[iy+2][ix][iz]+a3*p[iy+3][ix][iz]+
          a3*p[iy][ix-3][iz]+a2*p[iy][ix-2][iz]+a1*p[iy][ix-1][iz]+a0*p[iy][ix][iz]+
          a1*p[iy][ix+1][iz]+a2*p[iy][ix+2][iz]+a3*p[iy][ix+3][iz]+
          a3*p[iy][ix][iz-3]+a2*p[iy][ix][iz-2]+a1*p[iy][ix][iz-1]+a0*p[iy][ix][iz]+
          a1*p[iy][ix][iz+1]+a2*p[iy][ix][iz+2]+a3*p[iy][ix][iz+3]));
    }
  }
}
```


# Convolution



# Convolution

```
for(iy=3;iy < ny-4; ++iy){  
  for (ix=3; ix<nx-4; ++ix) {  
    for (iz=3; iz<nz-4; ++iz) {  
      pp[iy][ix][iz] = 2.0*p[iy][ix][iz]-pm[iy][ix][iz] +  
        scale*dvw[ix][iz]*(  
          a3*p[iy-3][ix][iz]+a2*p[iy-2][ix][iz]+a1*p[iy-1][ix][iz]+a0*p[iy][ix][iz]+  
          a1*p[iy+1][ix][iz]+a2*p[iy+2][ix][iz]+a3*p[iy+3][ix][iz]+  
          a3*p[iy][ix-3][iz]+a2*p[iy][ix-2][iz]+a1*p[iy][ix-1][iz]+a0*p[iy][ix][iz]+  
          a1*p[iy][ix+1][iz]+a2*p[iy][ix+2][iz]+a3*p[iy][ix+3][iz]+  
          a3*p[iy][ix][iz-3]+a2*p[iy][ix][iz-2]+a1*p[iy][ix][iz-1]+a0*p[iy][ix][iz]+  
          a1*p[iy][ix][iz+1]+a2*p[iy][ix][iz+2]+a3*p[iy][ix][iz+3]));  
    }  
  }  
}
```


# Convolution



# Convolution

```
for(iy=3;iy < ny-4; ++iy){
for (ix=3; ix<nx-4; ++ix) {
    for (iz=3; iz<nz-4; ++iz) {
        pp[iy][ix][iz] = 2.0*p[iy][ix][iz]-pm[iy][ix][iz] +
            scale*dvw[ix][iz]*(
                a3*p[iy-3][ix][iz]+a2*p[iy-2][ix][iz]+a1*p[iy-1][ix][iz]+a0*p[iy][ix][iz]+
                a1*p[iy+1][ix][iz]+a2*p[iy+2][ix][iz]+a3*p[iy+3][ix][iz]+
                a3*p[iy][ix-3][iz]+a2*p[iy][ix-2][iz]+a1*p[iy][ix-1][iz]+a0*p[iy][ix][iz]+
                a1*p[iy][ix+1][iz]+a2*p[iy][ix+2][iz]+a3*p[iy][ix+3][iz]+
                a3*p[iy][ix][iz-3]+a2*p[iy][ix][iz-2]+a1*p[iy][ix][iz-1]+a0*p[iy][ix][iz]+
                a1*p[iy][ix][iz+1]+a2*p[iy][ix][iz+2]+a3*p[iy][ix][iz+3]));
    }
}
}
```

# Convolution




# OpenMP Arrays

- Arrays are locally allocated
- Input arrays should generally be left unchanged
- Generally don't allocate/deallocate in parallel region



# CPU will allocate memory locally



# Depending on your NUMA machine

```
void main(int argc, char **argv){  
.  
.  
  
nth=1; //For serial case  
#pragma omp parallel  
nth=omp_get_num_threads();  
#pragma omp end parallel  
  
in=(float**)alloc_2d_float(n1,n2);  
work=(float**)alloc_2d_float(n1,nth);  
out=(float**) alloc_2d_float(n1,n2);  
  
#pragma omp parallel for  
for(i=0; i < n2; i++){  
    ith=omp_get_thread_num();  
    serial_work(in,out,work,i,ith);  
}
```

```
void serial_work(float **in, float **out, float **work, int i,  
int ith){  
    float alpha;  
    .  
    for(j=0; j < n1; j++){  
        .  
        .  
        alpha=func(in[i][j]....)  
        work[ith][j]=func(in[i][j].....)  
        .  
        .  
        out[i][j]=func(work[j]...)  
    }  
}
```

# Depending on your NUMA machine

```
void main(int argc, char **argv){  
.  
.  
  
nth=1; //For serial case  
#pragma omp parallel  
nth=omp_get_num_threads();  
#pragma omp end parallel  
  
in=(float**)alloc_2d_float(n1,n2);  
work=(float**)malloc(sizeof(float*)*nth);  
out=(float**) alloc_2d_float(n1,n2);  
  
#pragma omp parallel  
{  
    ith=omp_get_thread_num();  
    work[ith]=(float*)malloc(sizeof(float)*n1);  
    #pragma omp for private(i)  
    for(i=0; i < n2; i++){  
        serial_work(in,out,work,i,ith);  
    }  
}
```

```
void serial_work(float **in, float **out, float **work, int i,  
int ith){  
    float alpha;  
    .  
    .  
    for(j=0; j < n1; j++){  
        .  
        .  
        alpha=func(in[i][j]....)  
        work[ith][j]=func(in[i][j].....)  
        .  
        .  
        out[i][j]=func(work[j]...)  
    }  
}
```

Memory allocate locally

# OpenMP essentials

- Begin and end parallel block
- Private vs public variables
- Data parallelism with for/do
- Intermediate arrays expand dimensionality by nthreads

# OpenMP: What to remember when parallelizing

- Coarser grain the better
- Limit memory contention (most compute per word)
- Parallelism is often limited by memory bus

# Nested parallelism

```
program test_nested
integer, external :: omp_get_thread_num
integer :: main

write(0,*) "YEAH"
!$OMP PARALLEL private(main)
main=omp_get_thread_num()
write(0,*) "ONE",main
!$OMP PARALLEL
write(0,*) "TWO",omp_get_thread_num(),main
!$OMP END PARALLEL
!$OMP END PARALLEL

end program
```

# Nested parallelism

```
program test_nested
integer, external :: omp_get_thread_num
integer :: main

write(0,*) "YEAH"
!$OMP PARALLEL private(main)
main=omp_get_thread_num()
write(0,*) "ONE",main
!$OMP PARALLEL
write(0,*) "TWO",omp_get_thread_num(),main
!$OMP END PARALLEL
!$OMP END PARALLEL

end program
```

What is going to be  
the result?

# Nested parallelism: Default result

```
program test_nested
integer, external :: omp_get_thread_num
integer :: main

write(0,*) "YEAH"
!$OMP PARALLEL private(main)
main=omp_get_thread_num()
write(0,*) "ONE",main
!$OMP PARALLEL
write(0,*) "TWO",omp_get_thread_num(),main
!$OMP END PARALLEL
!$OMP END PARALLEL

end program
```

YEAH		
ONE	0	
TWO	0	0
ONE	2	
TWO	0	2
ONE	1	
ONE	3	
TWO	0	3
TWO	0	1



# Nested parallelism

setenv OMP\_NESTED TRUE


```
program test_nested
integer, external :: omp_get_thread_num
integer :: main

write(0,*) "YEAH"
!$OMP PARALLEL private(main)
main=omp_get_thread_num()
write(0,*) "ONE",main
!$OMP PARALLEL
write(0,*) "TWO",omp_get_thread_num(),main
!$OMP END PARALLEL
!$OMP END PARALLEL


end program
```

YEAH		
ONE	0	
TWO	0	0
TWO	3	0
ONE	1	
TWO	0	1
ONE	2	
TWO	1	0
ONE	3	
TWO	2	0
TWO	3	1
TWO	2	1
TWO	0	2
TWO	0	3
TWO	1	2
TWO	3	2
TWO	2	2
TWO	3	3
TWO	2	3
TWO	1	1
TWO	1	3

# Nested parallelism: I



# Nested parallelism: 2



# Thread parallelism

program thread_parallelism	CHECK THIS	88	3
integer :: i	CHECK THIS	89	3
integer, external :: omp_get_thread_num	CHECK THIS	90	3
real :: buf(100)	CHECK THIS	91	3
	CHECK THIS	92	3
!\$OMP PARALLEL	CHECK THIS	93	3
do i=1,100	CHECK THIS	94	3
!\$OMP TASK	CHECK THIS	95	3
buf(i)=i	CHECK THIS	96	3
write(0,*) "CHECK THIS",i,omp_get_thread_num()	CHECK THIS	98	3
!\$OMP END TASK	CHECK THIS	99	3
end do	CHECK THIS	97	0
!\$OMP END PARALLEL	CHECK THIS	86	1
	CHECK THIS	100	2
end program			

# Thread parallelism

- Create a series of tasks
- Tasks could be executed immediately or one resources are available
- Useful for
  - Consumer/producer
  - Recursive
  - Unbounded loops

# Task scoping

- shared
- private
- firstprivate - default for all non-shared variables
  - data is captured at creation
- default(shared|none)

# Thread parallelism

```
int a ;  
void foo ( ) {  
    int b , c ;  
    #pragma omp parallel shared ( b )  
    #pragma omp parallel private ( b )  
    {  
        int d ;  
        #pragma omp task  
        {  
            int e ;  
            a =  
            b =  
            c =  
            d =  
            e =  
        }  
    }  
}
```

# Thread parallelism

```
int a ;
void foo ( ) {
  int b , c ;
  #pragma omp parallel shared ( b )
  #pragma omp parallel private ( b )
  {
    int d ;
    #pragma omp task
    {
      int e ;
      a = shared
      b =
      c =
      d =
      e =
    }
  }
}
```



# Thread parallelism

```
int a ;  
void foo ( ) {  
    int b , c ;  
    #pragma omp parallel shared ( b )  
    #pragma omp parallel private ( b )  
    {  
        int d ;  
        #pragma omp task  
        {  
            int e ;  
            a = shared  
            b = firstprivate  
            c =  
            d =  
            e =  
        }  
    }  
}
```

# Thread parallelism

```
int a ;  
void foo ( ) {  
    int b , c ;  
    #pragma omp parallel shared ( b )  
    #pragma omp parallel private ( b )  
    {  
        int d ;  
        #pragma omp task  
        {  
            int e ;  
            a = shared  
            b = firstprivate  
            c = shared  
            d =  
            e =  
        }  
    }  
}
```

# Thread parallelism

```
int a ;  
void foo ( ) {  
    int b , c ;  
    #pragma omp parallel shared ( b )  
    #pragma omp parallel private ( b )  
    {  
        int d ;  
        #pragma omp task  
        {  
            int e ;  
            a = shared  
            b = firstprivate  
            c = shared  
            d = firstprivate  
            e =  
        }  
    }  
}
```

# Thread parallelism

```
int a ;  
void foo ( ) {  
    int b , c ;  
    #pragma omp parallel shared ( b )  
    #pragma omp parallel private ( b )  
    {  
        int d ;  
        #pragma omp task  
        {  
            int e ;  
            a = shared  
            b = firstprivate  
            c = shared  
            d = firstprivate  
            e = private  
        }  
    }  
}
```

# OpenMP tasks

- Barrier

`!$OMPTASKWAIT`

- if statements

# Nested parallelism

setenv OMP\_NESTED TRUE

```
program nested_two
integer, external :: omp_get_num_threads
integer, external :: omp_get_thread_num
integer :: main, nlocal, ntot

ntot=omp_get_num_threads()
write(0,*) "YEAH"
call omp_set_num_threads(2)
!$OMP PARALLEL private(main)
main=omp_get_thread_num()
call omp_set_num_threads(ntot/2)
write(0,*) "ONE",main
!$OMP PARALLEL
write(0,*) "TWO",omp_get_thread_num(),main
!$OMP END PARALLEL
!$OMP END PARALLEL
end program
```

YEAH		
ONE	0	
TWO	0	0
ONE	1	
TWO	0	1