

Roofline model

Serial machines

What is the maximum performance possible for a given algorithm?

- What is your bottleneck?
 - Disk?
 - Math ops?
 - Reading from cache?
 - Reading from main memory?

First a little on counting reads and ops

```
for(i2=0; i2 < n2; i2++){  
    out[i2][1] += in[i2][0] * val[i2]  
}
```

First a little on counting reads and ops

Reads: 5

Writes: 1

```
for(i2=0; i2 < n2; i2++){  
    out[i2][1] += in[i2][0] * val[i2]  
}
```

First a little on counting reads and ops

Reads: 5

Writes: 1

```
for(i2=0; i2 < n2; i2++){  
    out[i2][1] += in[i2][0] * val[i2]  
}
```

Multiply: 1

Add: 1

First a little on counting reads and ops

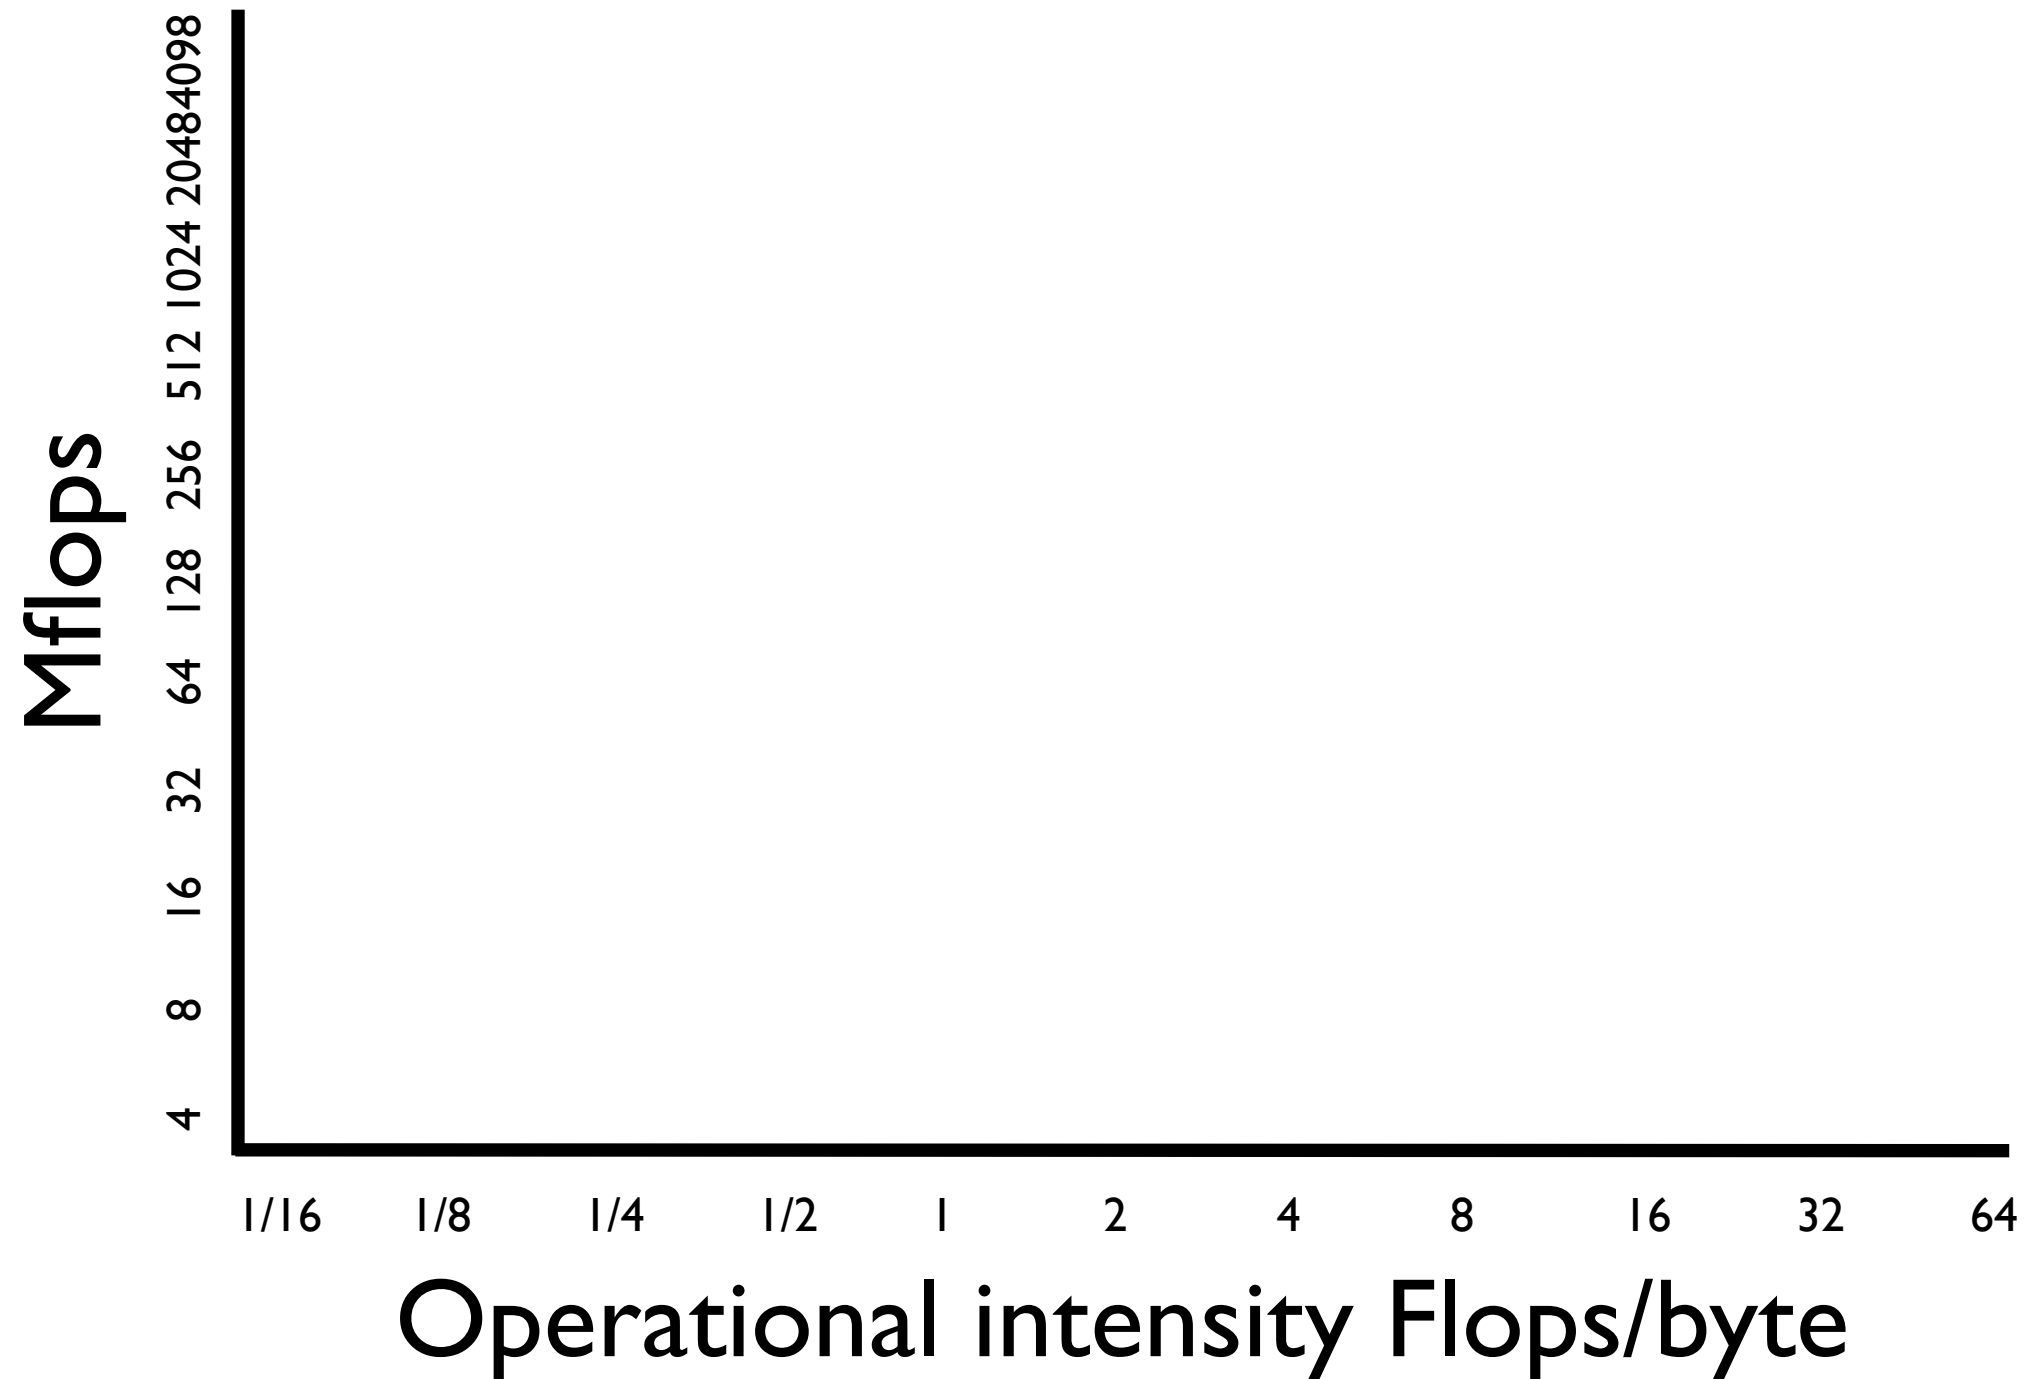
Reads: 5
Writes: 1

```
for(i2=0; i2 < n2; i2++){  
    out[i2][1] += in[i2][0] * val[i2]  
}
```

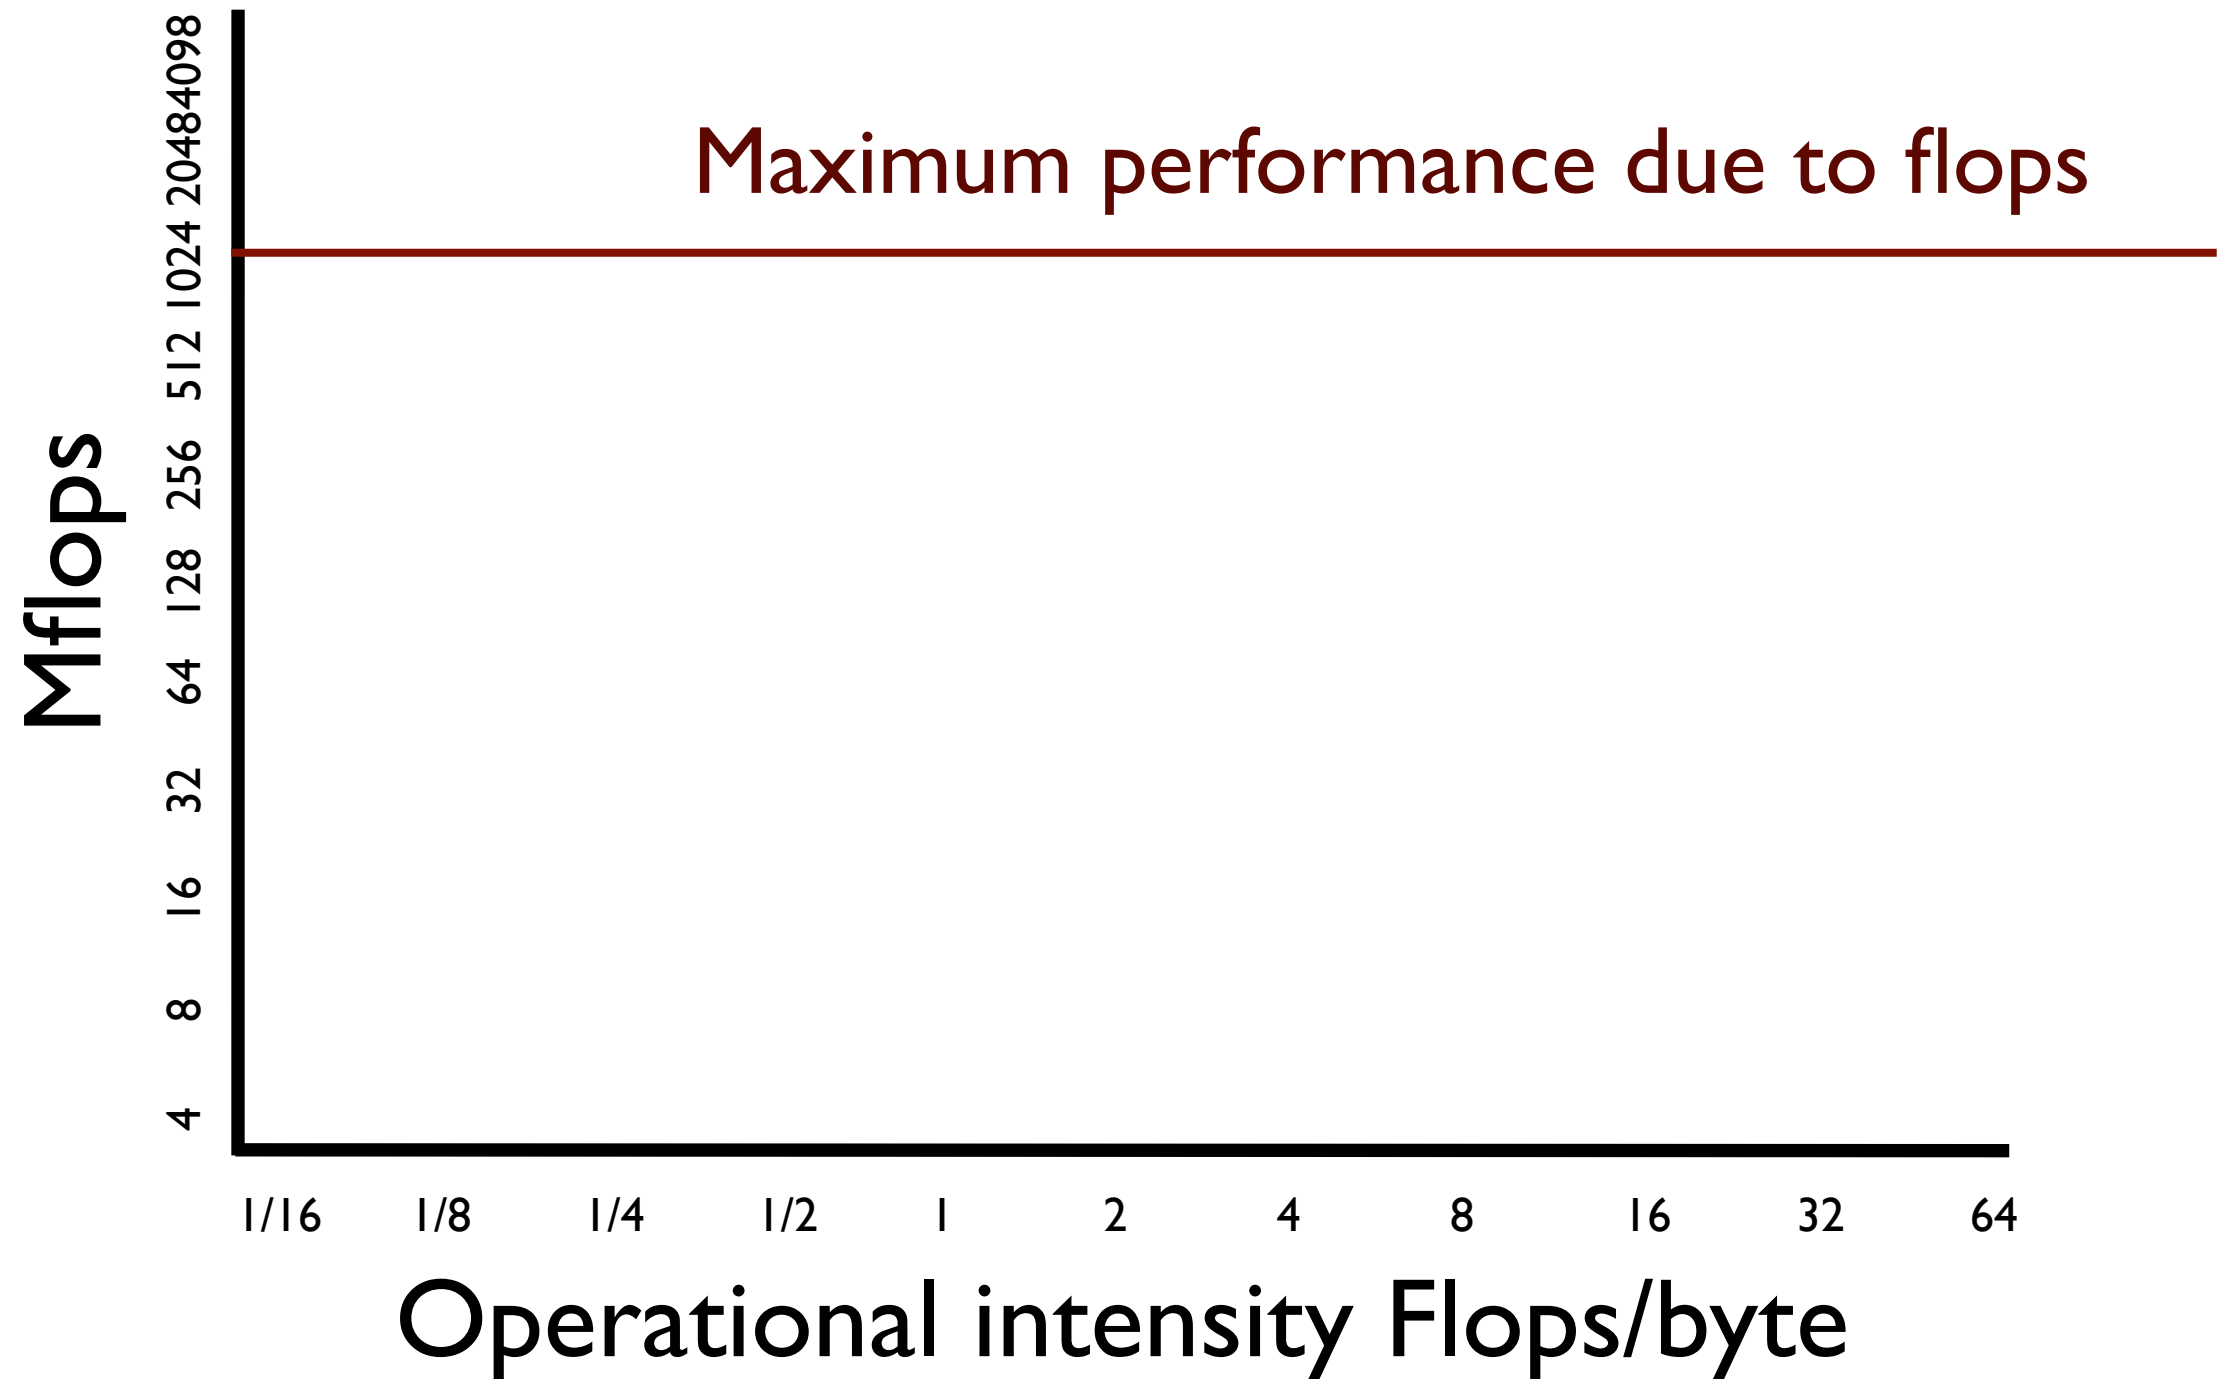
Multiply: 1
Add: 1

IO: 24
Flops: 2
Ratio: 1/12

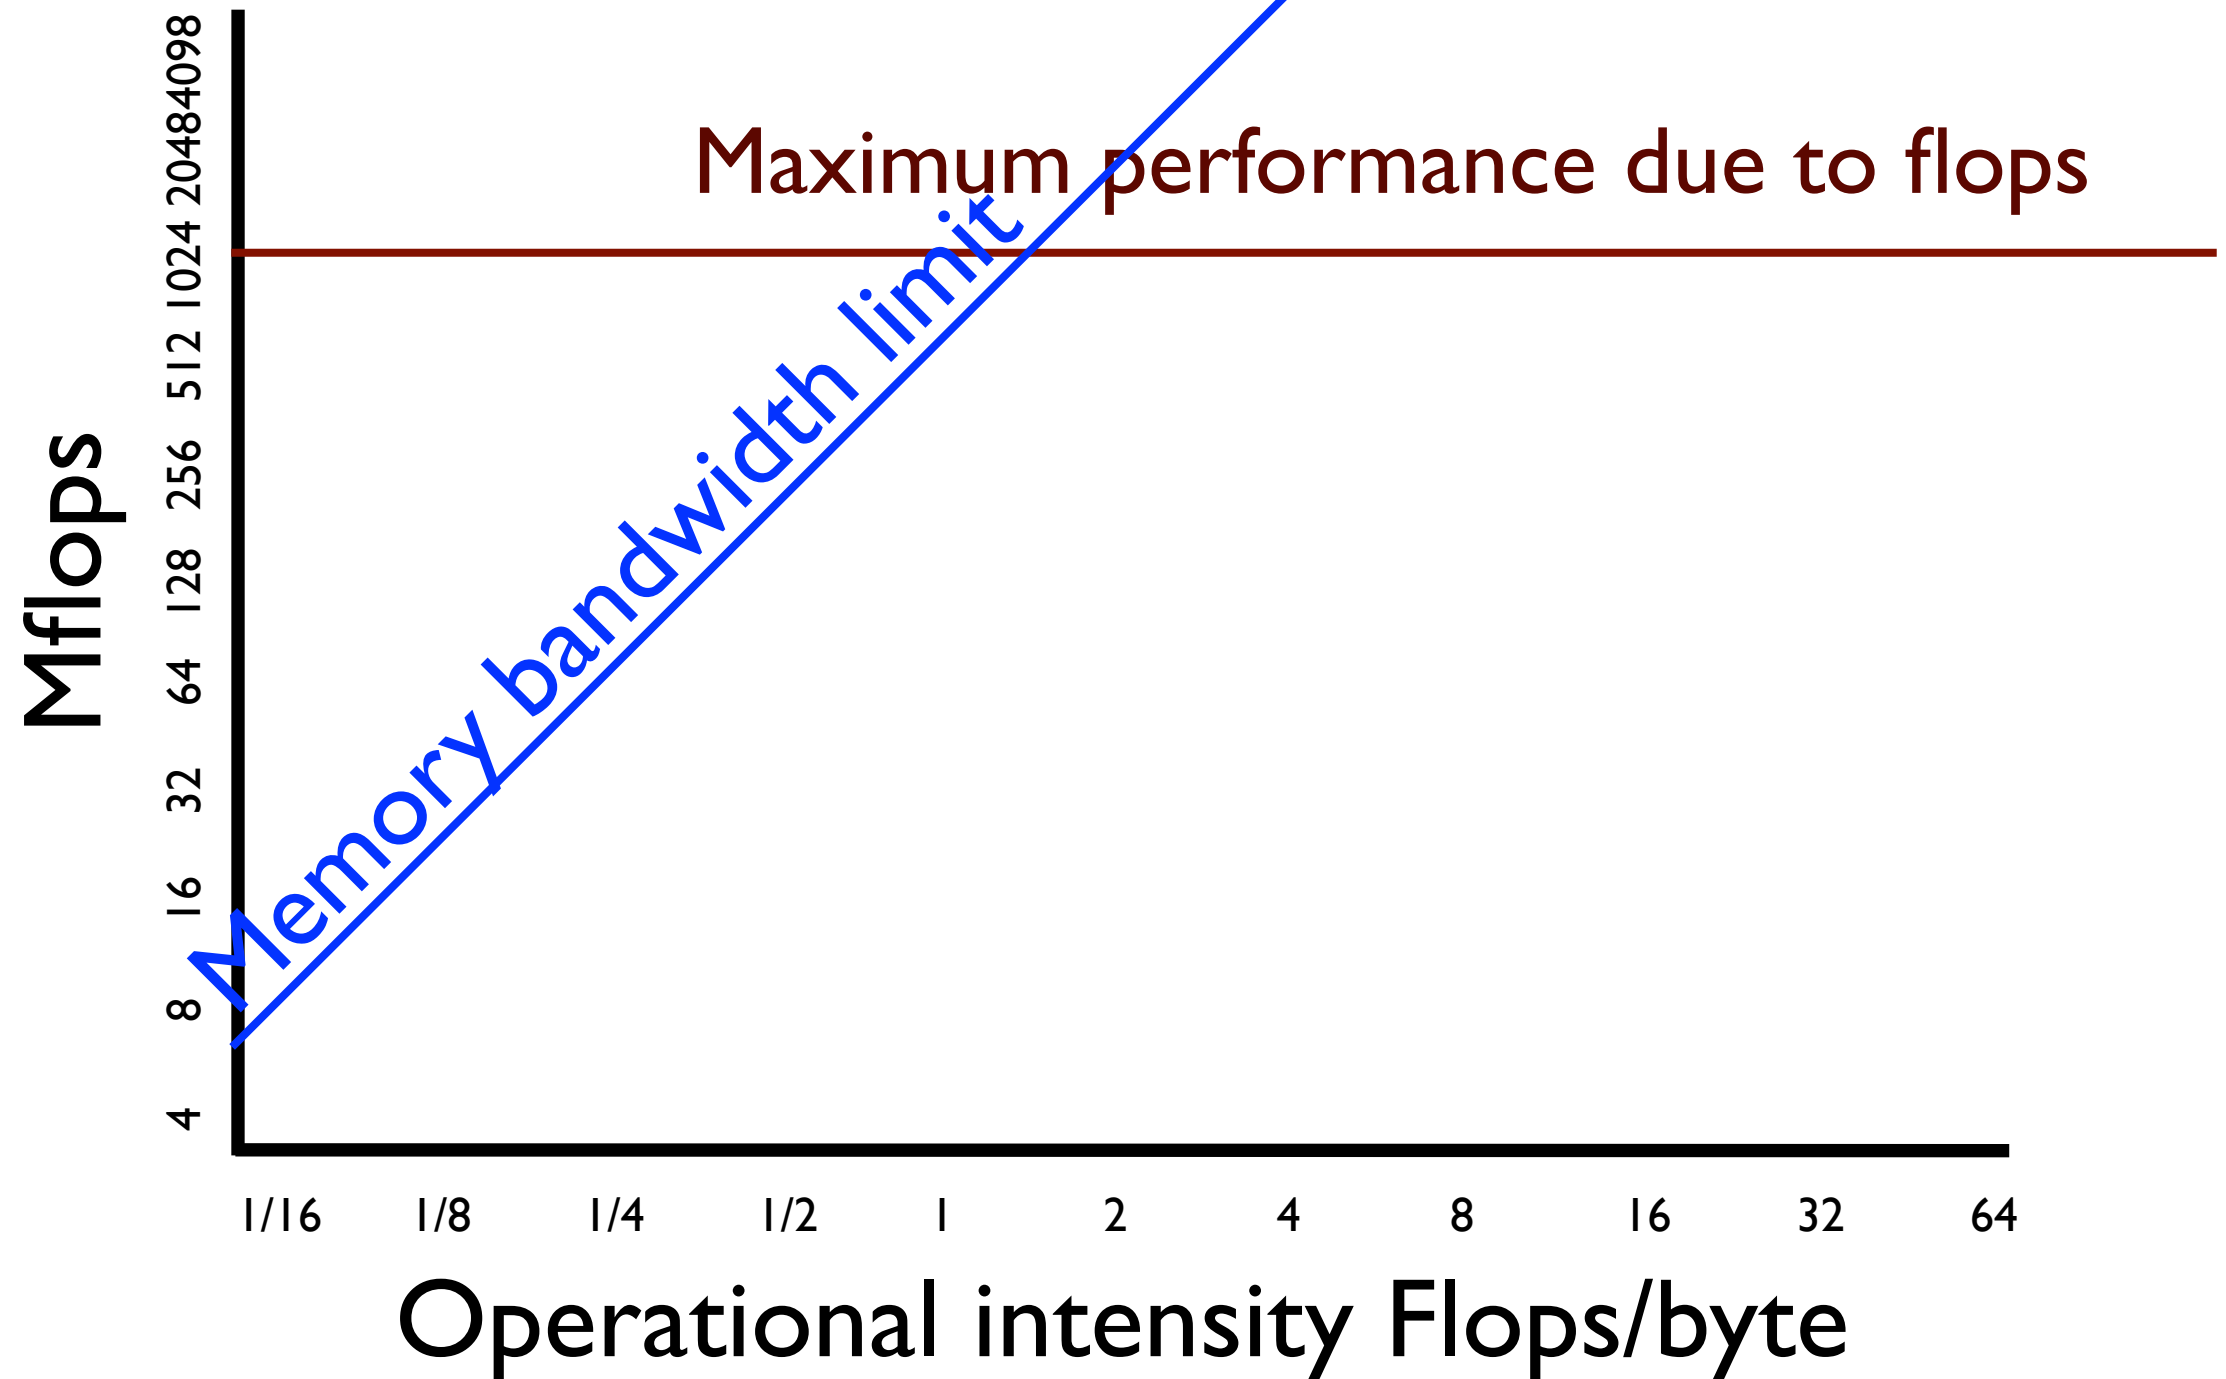
Roofline model




Roofline model



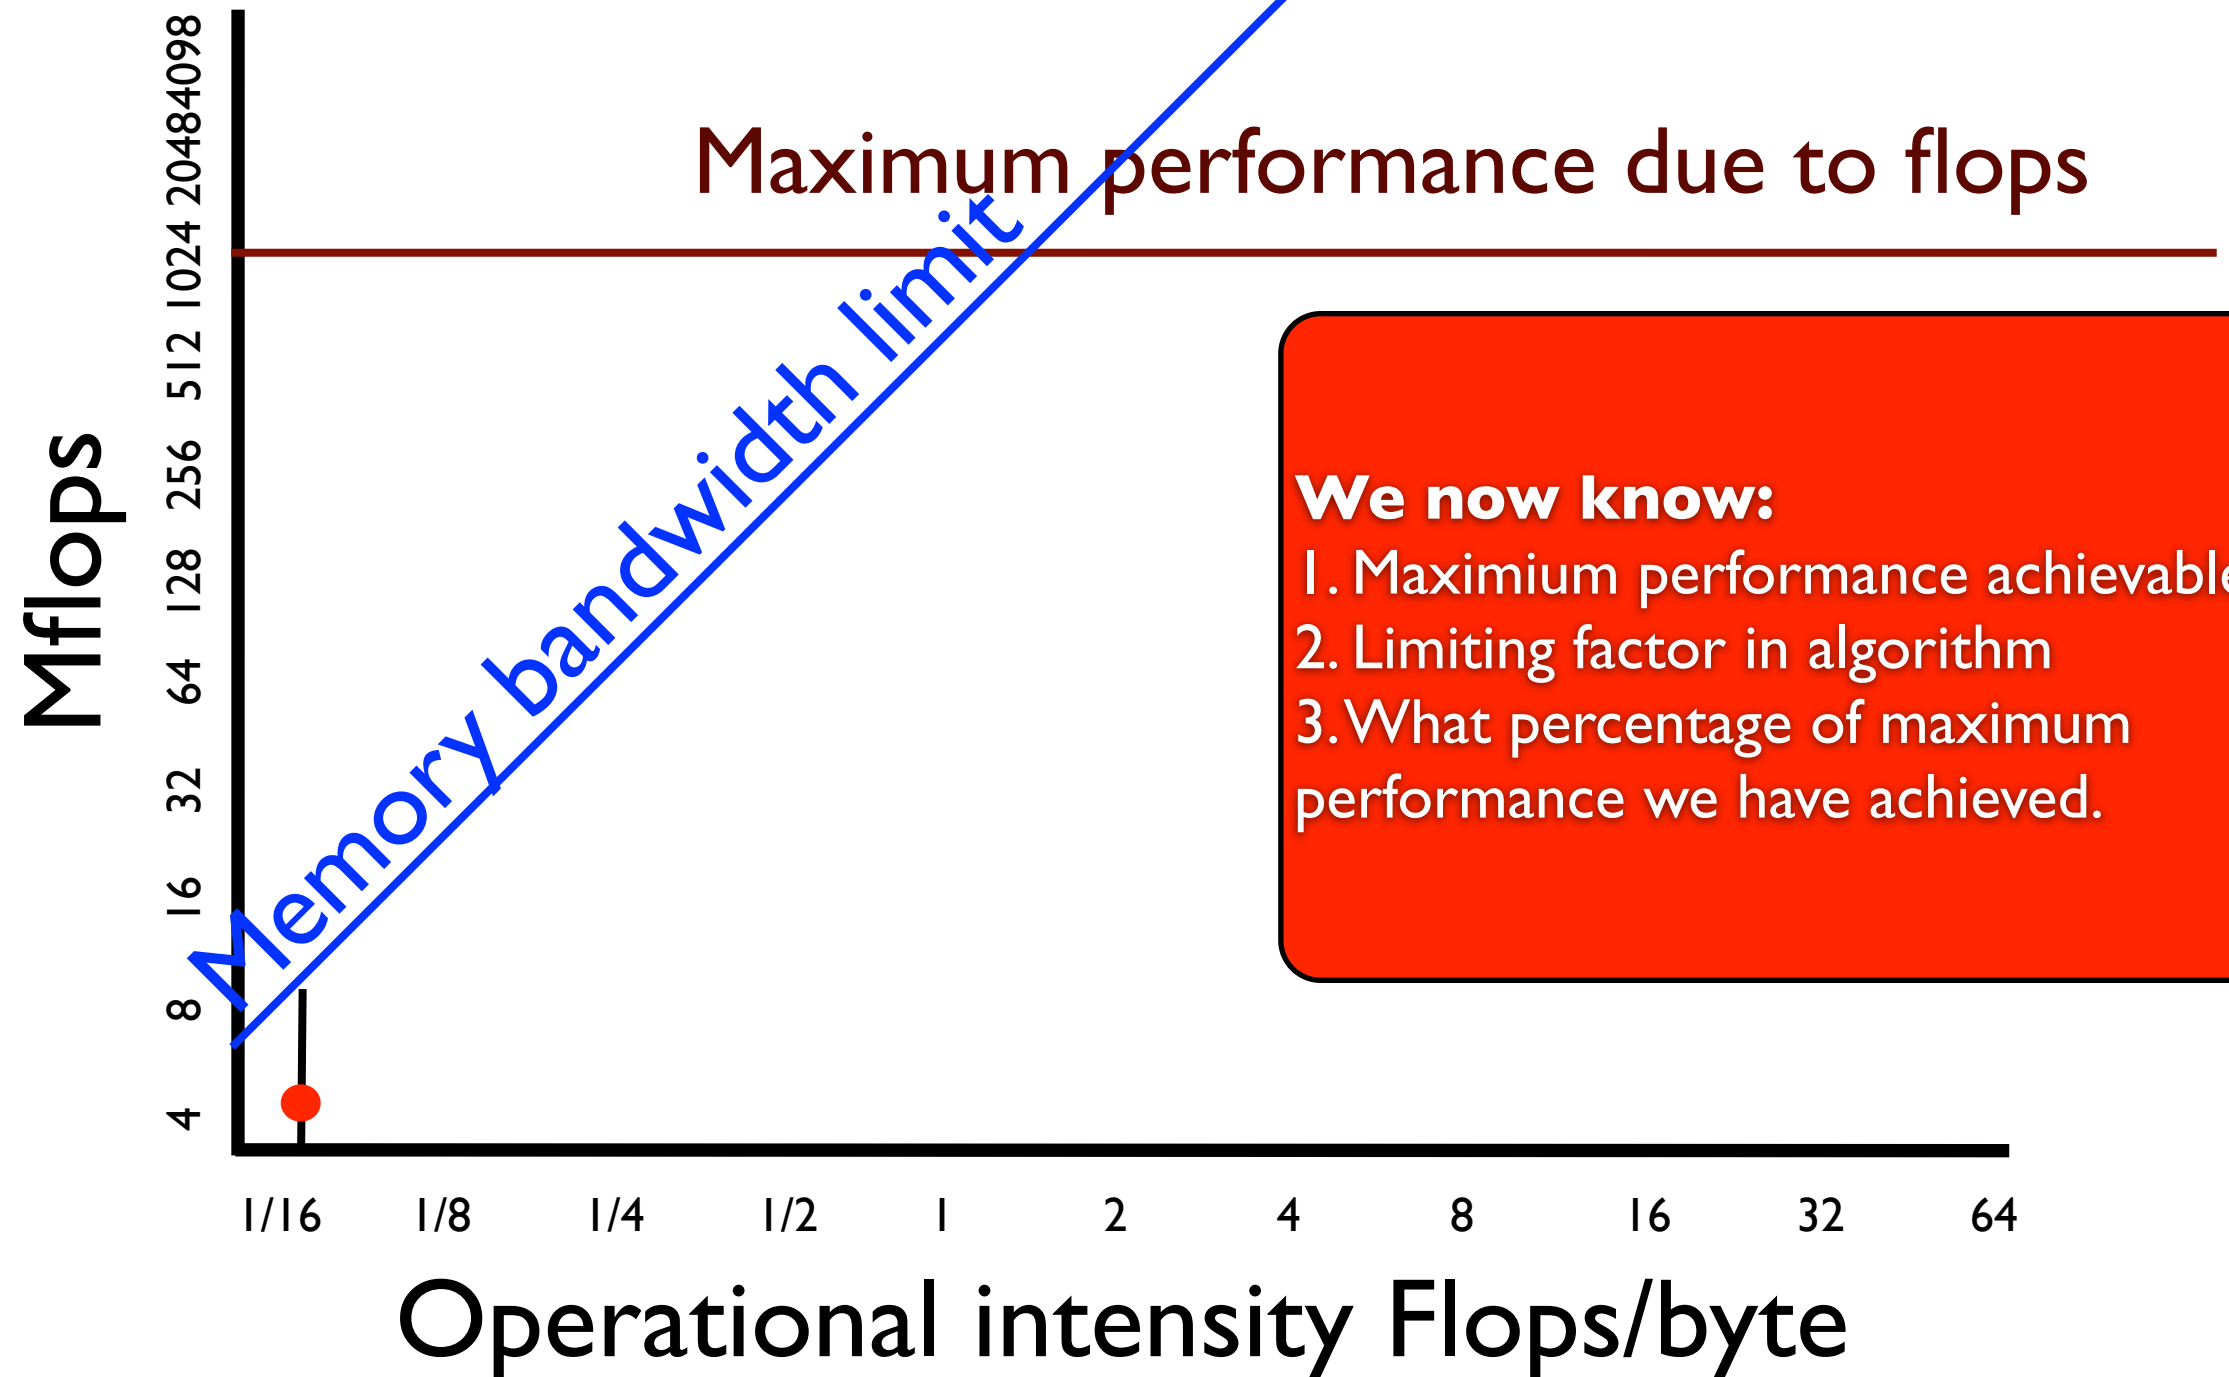
Roofline model



Roofline model



Roofline model



We now know:

1. Maximum performance achievable
2. Limiting factor in algorithm
3. What percentage of maximum performance we have achieved.

Causes for not achieving maximum performance

1. Indirection (iloc needs to be looked up before knowing what to grab from in, out)
2. Poor use of cache line for in,out
3. Inability to use prefetch on in,out

```
for(i2=0; i2 < n2; i2++){  
    out[iloc[i2][1]]+=in[iloc[i2][0]] *val[i2]  
}
```

What if the relative filter locations were
always the same?

```
i=0;
for(i2=3; i2 < n2-3; i2++){
    ib=i2-3;
    for(i1=0; i1 < 7; i1++){
        out[i2]+=in[ib+i1]*val[i1];
    }
}
```

What if the relative filter locations were always the same?

```
i=0;
```

```
for(i2=3; i2 < n2-3; i2++){
```

```
    ib=i2-3;
```

```
    for(i1=0; i1 < 7; i1++){
```

```
        out[i2]+=in[ib++]*val[i++];
```

```
    }
```

Reads: 3

Writes: 1

Multiply: 1

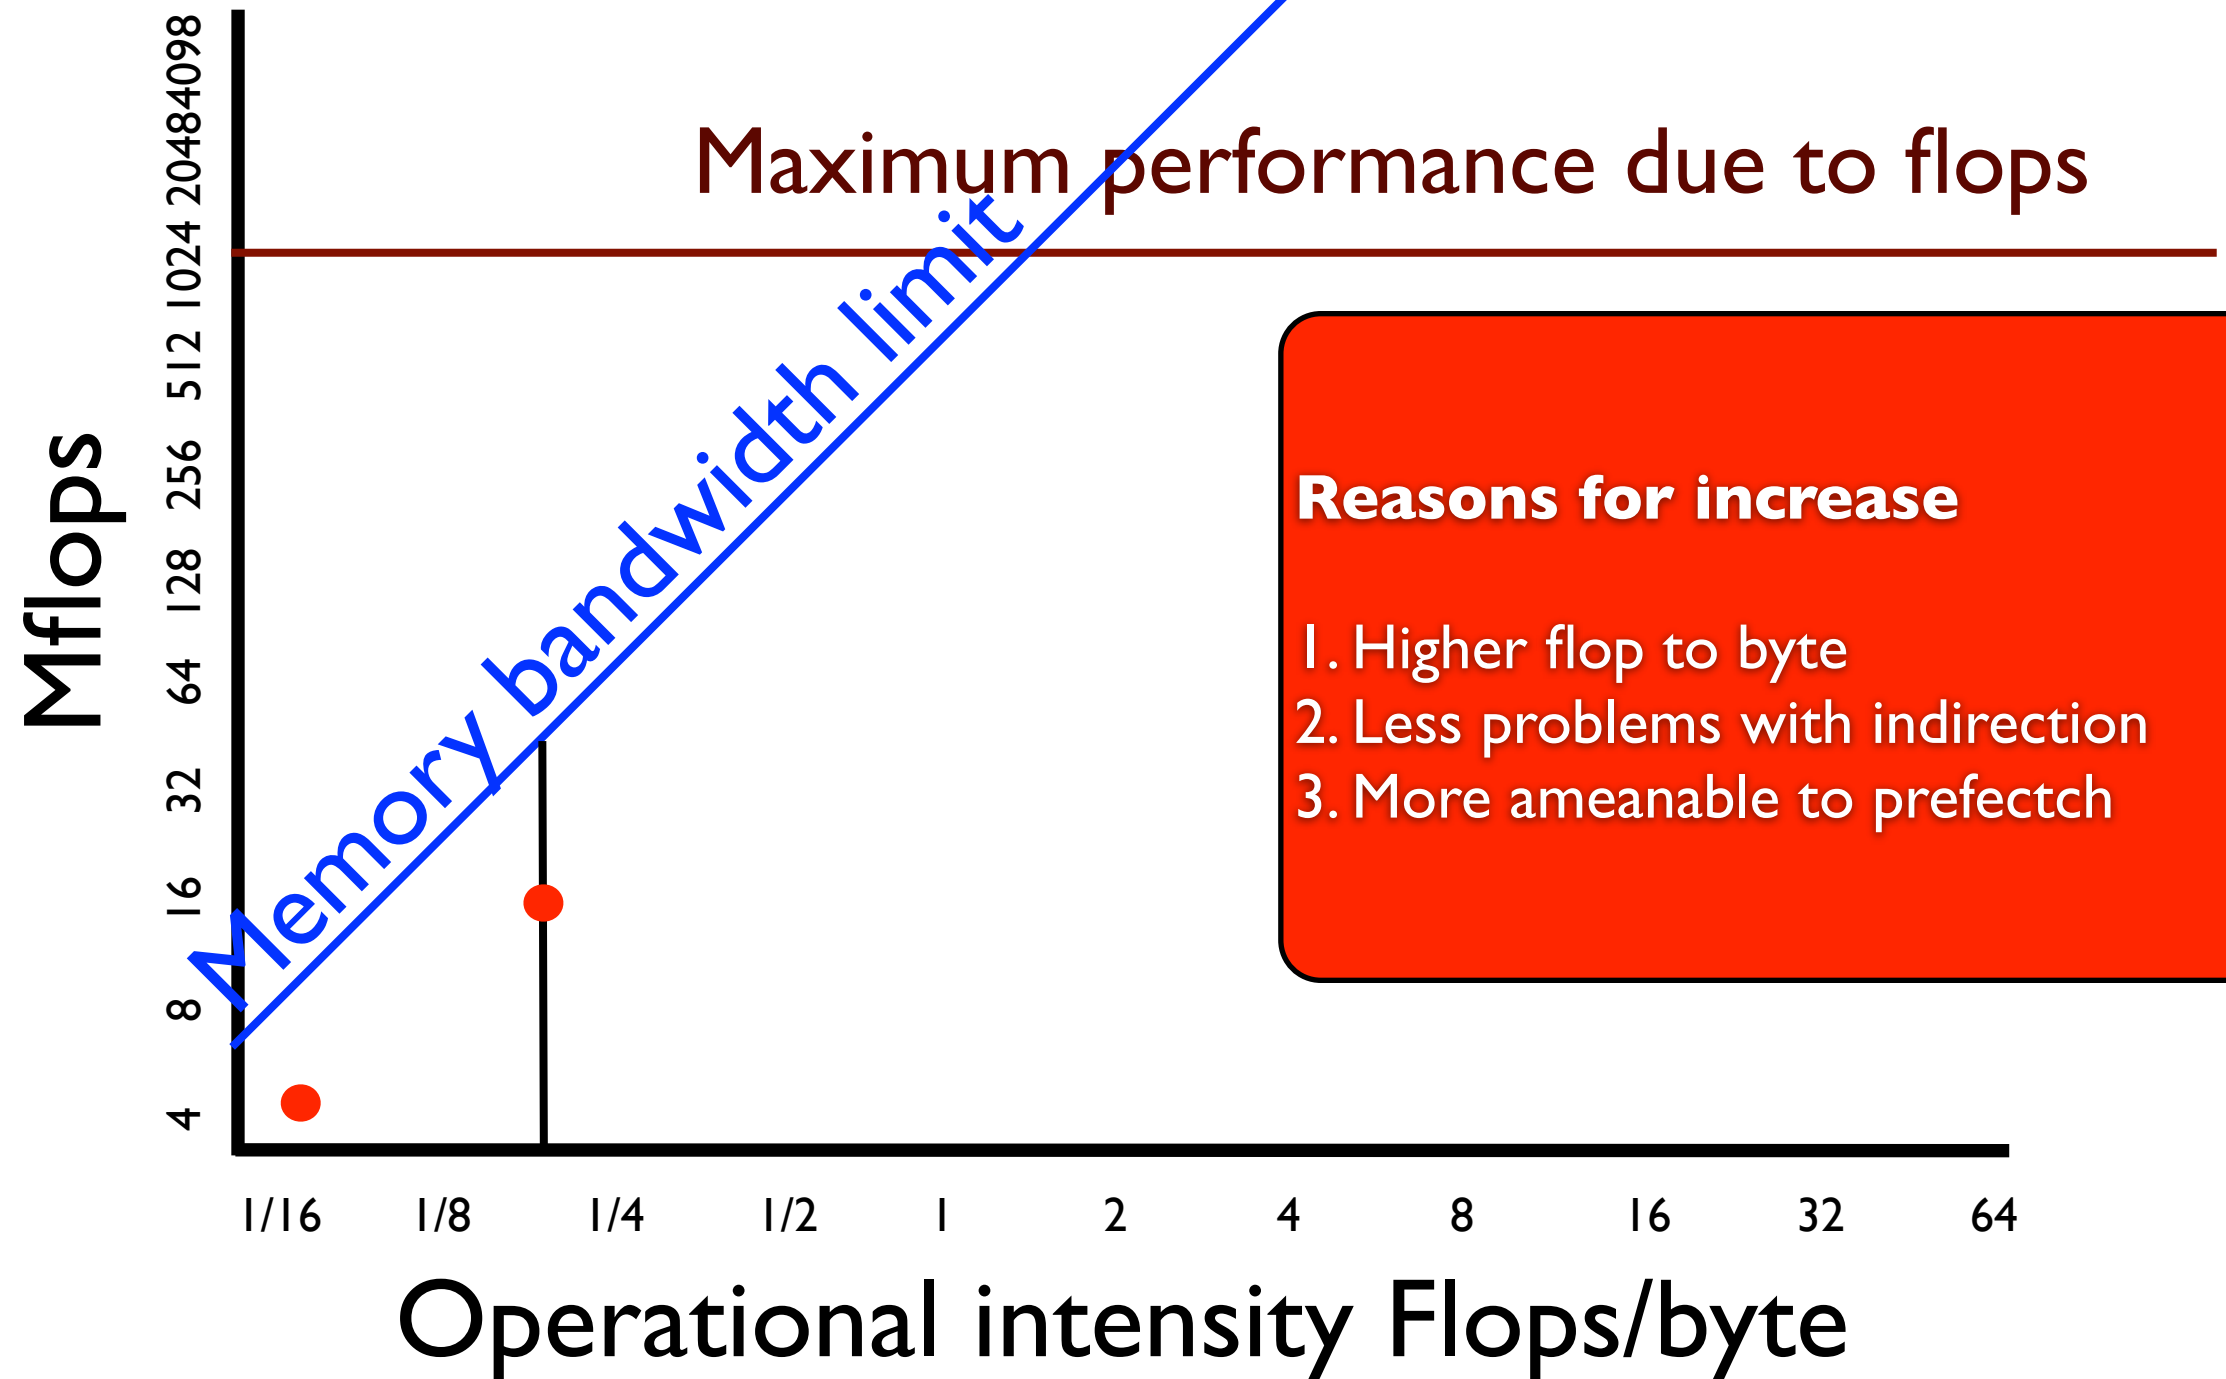
Add: 1

IO: 12

Flops: 2

Ratio: 1/6

Roofline model



But we aren't dealing with a Von
Neumann machine, we have cache

```
i=0;  
for(i2=3; i2 < n2-3; i2++){  
    ib=i2-3;  
    for(i1=0; i1 < 7; i1++){  
        out[i2]+=in[ib++]*val[i++];  
    }  
}
```

In

Has a reuse rate of 7. Only when $i1=6$ does a new value need to be read in from memory

out

Has a reuse rate of $1/7$. Only when $i1=0$ does a new value need to be read. Only when $i1=6$ does a new value need to be written.

IO: 9.1


Flops: 2

Ratio: $1/4.5$

Reads: $11/7$

Writes: $1/7$

Roofline model



What if val is only a function of i1?

```
for(i2=3; i2 < n2-3; i2++){  
    ib=i2-3;i=0;  
    for(i1=0; i1 < 7; i1++){  
        out[i2]+=in[ib++]*val[i1++];  
    }  
}
```

What if val is only a function of i1?

```
for(i2=3; i2 < n2-3; i2++){  
    ib=i2-3;i=0;  
    for(i1=0; i1 < 7; i1++){  
        out[i2]+=in[ib++]*val[i++];  
    }  
}
```

IO: 1.12

Flops:2

Ratio: 1/4.5

Reads: 1/7

Writes: 1/7

Roofline model

